

Next-generation Bayesian local robustness

Ryan Giordano (rgiordano@berkeley.edu, UC Berkeley)

Berkeley Neyman Seminar


September 2026

The basic problem


We have a survey population, for whom we observe:

- Covariates x (e.g. race, gender, zip code, age, education level)
- Responses y (e.g. A binary response to “do you support candidate Z”)

We want the average response in a target population, in which we observe only covariates.



Observe (x_i, y_i) for $i = 1, \dots, N_S$



Observe x_j for $j = 1, \dots, N_T$

The basic problem

We have a survey population, for whom we observe:

- Covariates x (e.g. race, gender, zip code, age, education level)
- Responses y (e.g. A binary response to “do you support candidate Z”)


We want the average response in a target population, in which we observe only covariates.

The problem is that the populations may be very different, maybe leading to bias.



Photo copyright Mark Taylor: natureimg.com

Observe (x_i, y_i) for $i = 1, \dots, N_S$



Observe x_j for $j = 1, \dots, N_T$

The basic problem

We have a survey population, for whom we observe:

- Covariates x (e.g. race, gender, zip code, age, education level)
- Responses y (e.g. A binary response to “do you support candidate Z”)


We want the average response in a target population, in which we observe only covariates.

The problem is that the populations may be very different, maybe leading to bias.



Photo copyright Mark Taylor: naturepl.com

Observe (x_i, y_i) for $i = 1, \dots, N_S$



Observe x_j for $j = 1, \dots, N_T$

This talk will study “Multilevel regression and poststratification” estimators:

- Set $\mathbb{E}[y|x, \theta] = m(\theta^\top x)$ and prior $\mathcal{P}(\theta)$ (e.g. Hierarchical logistic regression)
- Estimate the posterior $\mathcal{P}(\theta|\text{Data})$ with Markov Chain Monte Carlo (MCMC)
- Take $\hat{\mu}^{\text{MrP}}(Y_S) = \frac{1}{N_T} \sum_{j=1}^{N_T} \hat{y}_j$ where $\hat{y}_j = \mathbb{E}_{\mathcal{P}(\theta|\text{Data})}[y|x_j]$.

The basic problem

We have a survey population, for whom we observe:

- Covariates x (e.g. race, gender, zip code, age, education level)
- Responses y (e.g. A binary response to “do you support candidate Z”)


We want the average response in a target population, in which we observe only covariates.

The problem is that the populations may be very different, maybe leading to bias.



Photo copyright Mark Taylor: naturepl.com

Observe (x_i, y_i) for $i = 1, \dots, N_S$



Observe x_j for $j = 1, \dots, N_T$

This talk will study “Multilevel regression and poststratification” estimators:

- Set $\mathbb{E}[y|x, \theta] = m(\theta^\top x)$ and prior $\mathcal{P}(\theta)$ (e.g. Hierarchical logistic regression)
- Estimate the posterior $\mathcal{P}(\theta|\text{Data})$ with Markov Chain Monte Carlo (MCMC)
- Take $\hat{\mu}^{\text{MrP}}(Y_S) = \frac{1}{N_T} \sum_{j=1}^{N_T} \hat{y}_j$ where $\hat{y}_j = \mathbb{E}_{\mathcal{P}(\theta|\text{Data})}[y|x_j]$.

Problem: MCMC is time consuming and the regression model is surely misspecified.

Goal: Meaningful regression specification checks with frequentist uncertainty.

Result preview




Figure 1: A covariate shift diagnostic

Here is our check for two models on the same simulated data: one correctly specified and one misspecified. The misspecified model is more extreme.

Problem: What is “large enough” to be real evidence of misspecification?

Possible answer: Frequentist confidence intervals via the parametric bootstrap.

Result preview




Figure 1: A covariate shift diagnostic

The frequentist interval based on the parametric bootstrap rejects the misspecified model but not the correctly specified model.

Problem: Expensive: each bootstrap sample required a new MCMC run.

Possible answer: Infinitesimal jackknife confidence intervals? [Giordano and Broderick, 2026]

Result preview




Figure 1: A covariate shift diagnostic

IJ confidence intervals are fast and often work well for posterior means.
But they are not accurate for this more complex diagnostic (high Monte Carlo error).

Result preview




Figure 1: A covariate shift diagnostic

IJ confidence intervals are fast and often work well for posterior means. But they are not accurate for this more complex diagnostic (high Monte Carlo error).

Answer: Our flow-based local sensitivity measures (“Otto grad”) produce good approximations to the bootstrap at negligible computational expense.

- Bayesian local regression specification checks
 - (With Alice Cima, Erin Hartman, Jared Miller, Avi Feller)
- Frequentist variability of Bayesian posteriors with the infinitesimal jackknife
 - (With Tamara Broderick)
- Flow-based Bayesian local robustness: “Otto grad”
 - (With Lucas Schwengber)

Bayesian local regression specification checks

Local specification checks

$$\hat{\mu}^{\text{MrP}}(Y_S) = \frac{1}{N_T} \sum_{j=1}^{N_T} \hat{y}_j \text{ where } \hat{y}_j = \underbrace{\mathbb{E}_{\mathcal{P}(\theta|\text{Data})} [m(\theta^\top \mathbf{x}_j)]}_{\text{Want to check without re-running MCMC}}.$$

This model of *survey data* is definitely misspecified.

Why do we care about correct specification? We want to use $\mathcal{P}_S(x)$ to learn about $\mathcal{P}_T(x)$.

Correct specification $\Rightarrow$ Regression estimates are invariant to $\mathcal{P}_S(x)$

Idea: Check invariance to $\mathcal{P}_S(x)$ directly.

Let $\mathcal{P}(x, y)$ denote a generic joint density of x, y .

Let $\hat{\mathcal{P}}_S$ the empirical distribution on the survey data.

1. Write our estimand as a functional $\mu(\mathcal{P})$ that takes in a density and returns a MrP estimate
2. Define a covariate density perturbation $w(x)$
3. Define $\mathcal{P}_\epsilon(x, y) = \mathcal{P}(x, y) + \epsilon w(x) \mathcal{P}(y|x) = (\mathcal{P}(x) + \epsilon w(x)) \mathcal{P}(y|x)$
4. Compute the derivative $\Delta(w, \mathcal{P}) = \partial \mu(\mathcal{P}_\epsilon) / \partial \epsilon$
5. Plug the empirical distribution into $\Delta(w, \hat{\mathcal{P}}_S)$ and check whether ≈ 0

Bayesian sensitivity analysis

Consider a family of distributions $\mathcal{P}^\epsilon(\theta) \propto \mathcal{P}(\theta) \exp(\epsilon h(\theta))$ and a test function $\phi(\theta)$. Then, by interchange of integration and differentiation,

$$\left. \frac{\partial \mathbb{E}_{\mathcal{P}^\epsilon(\theta)} [\phi(\theta)]}{\partial \epsilon} \right|_{\epsilon=0} = \text{Cov}_{\mathcal{P}(\theta)} (\phi(\theta), h(\theta)) .$$

Let the regression log likelihood be $\ell(y|x, \theta)$, so that

$$\mathcal{P}(\theta|\text{Data}) \propto \exp \left(\sum_{i=1}^{N_S} \ell(y_i|x_i, \theta) \right) \text{ and } \mathcal{P}^\epsilon(\theta|\text{Data}) \propto \mathcal{P}(\theta|\text{Data}) \exp \left(\sum_{i=1}^{N_S} \epsilon w(x_i) \ell(y_i|x_i, \theta) \right)$$

- $\mathcal{P}(\theta) = \mathcal{P}(\theta|\text{Data})$
- $\phi(\theta) = \frac{1}{N_T} \sum_{j=1}^{N_T} m(\theta^\top x_j)$ (so that $\hat{\mu}^{\text{MrP}}(Y_S) = \mathbb{E}_{\mathcal{P}(\theta|\text{Data})} [\phi(\theta)]$)
- $h(\theta) = \sum_{i=1}^{N_S} w(x_i) \ell(y_i|x_i, \theta) \Rightarrow$

Our diagnostic is then given by

$$\Delta(w, \hat{\mathcal{P}}_S) := \left. \frac{\partial \hat{\mu}^{\text{MrP}}(Y_S)}{\partial \epsilon} \right|_{\epsilon=0} = \text{Cov}_{\mathcal{P}(\theta|\text{Data})} \left(\frac{1}{N_T} \sum_{j=1}^{N_T} m(\theta^\top x_j), \sum_{i=1}^{N_S} w(x_i) \ell(y_i|x_i, \theta) \right)$$

We expect $\Delta(w, \hat{\mathcal{P}}_S)$ to be close to zero when the model is correctly specified and posterior estimation is accurate.




Figure 2: A covariate shift diagnostic

Simulation data: Logistic regression with five active regressors, each with covariate shift.

Model 0 (correct) : $y \sim \text{Logistic}(1 + X_1 + X_2 + X_3 + X_4 + X_5)$

Model 1 (misspecified) : $y \sim \text{Logistic}(1 + X_3 + X_4 + X_5)$

Compute diagnostic with $w(x) = \mathbb{I}(X_2 > 0.878)$

Frequentist variance with the infinitesimal
jackknife

Why frequentist variance?

For the remainder of the talk, focus on approximating the *frequentist* sampling distribution under $x_i, y_i \stackrel{iid}{\sim} \mathcal{P}_S(x, y)$ of quantities like our diagnostic:

$$\Delta = \text{Cov}_{\mathcal{P}(\theta|\text{Data})} \left(\frac{1}{\textcolor{red}{N}_T} \sum_{j=1}^{\textcolor{red}{N}_T} m(\theta^\top x_j), \sum_{i=1}^{\textcolor{blue}{N}_S} w(x_i) \ell(y_i | x_i, \theta) \right)$$

No need to commit to our covariate shift diagnostic to care!

- Frequentist variability of *any* model checking diagnostic is interesting
- Frequentist variability remains meaningful under model misspecification (Huggins and Miller, IJ papers)

We will form *local approximations* computable only from *the original posterior*.

We will compare with computationally expensive bootstrap:

- Compute Δ on the original dataset
- For $B = 1, \dots, 100$:
 1. Sample $(x_1^*, y_1^*), \dots, (x_{\textcolor{blue}{N}_S}^*, y_{\textcolor{blue}{N}_S}^*)$ (parametric or nonparametric bootstrap)
 2. Run MCMC to estimate $\mathcal{P}(\theta | (x_1^*, y_1^*), \dots, (x_{\textcolor{blue}{N}_S}^*, y_{\textcolor{blue}{N}_S}^*))$
 3. Use the draws to estimate Δ_B^*
- Compare Δ with quantiles of $\Delta_B^* - \Delta$

Data re-weighting.

Augment the problem with *data weights* $w_1, \dots, w_N$. We can write $\mathbb{E}_{\mathcal{P}(\theta|X;w)} [f(\theta)]$.

$$\ell_n(\theta) := \log \mathcal{P}(x_n|\theta) \qquad \log \mathcal{P}(X|\theta; w) = \sum_{n=1}^N w_n \ell_n(\theta)$$

Data re-weighting.

Augment the problem with *data weights* $w_1, \dots, w_N$. We can write $\mathbb{E}_{\mathcal{P}(\theta|X;w)} [f(\theta)]$.

$$\ell_n(\theta) := \log \mathcal{P}(x_n|\theta) \qquad \log \mathcal{P}(X|\theta; w) = \sum_{n=1}^N w_n \ell_n(\theta)$$

Original weights:




Data re-weighting.

Augment the problem with *data weights* $w_1, \dots, w_N$. We can write $\mathbb{E}_{\mathcal{P}(\theta|X;w)} [f(\theta)]$.

$$\ell_n(\theta) := \log \mathcal{P}(x_n|\theta) \qquad \log \mathcal{P}(X|\theta; w) = \sum_{n=1}^N w_n \ell_n(\theta)$$

Original weights:



Leave-one-out weights:



Data re-weighting.

Augment the problem with *data weights* $w_1, \dots, w_N$. We can write $\mathbb{E}_{\mathcal{P}(\theta|X;w)} [f(\theta)]$.

$$\ell_n(\theta) := \log \mathcal{P}(x_n|\theta)$$

$$\log \mathcal{P}(X|\theta; w) = \sum_{n=1}^N w_n \ell_n(\theta)$$


Original weights:



Leave-one-out weights:



Bootstrap weights:




Data re-weighting.

Augment the problem with *data weights* $w_1, \dots, w_N$. We can write $\mathbb{E}_{\mathcal{P}(\theta|X;w)} [f(\theta)]$.


$$\ell_n(\theta) := \log \mathcal{P}(x_n|\theta)$$

$$\log \mathcal{P}(X|\theta; w) = \sum_{n=1}^N w_n \ell_n(\theta)$$



Original weights:



Leave-one-out weights:



Bootstrap weights:



Data re-weighting.

Augment the problem with *data weights* $w_1, \dots, w_N$. We can write $\mathbb{E}_{\mathcal{P}(\theta|X;w)} [f(\theta)]$.


$$\ell_n(\theta) := \log \mathcal{P}(x_n|\theta)$$

$$\log \mathcal{P}(X|\theta; w) = \sum_{n=1}^N w_n \ell_n(\theta)$$



Original weights:



Leave-one-out weights:



Bootstrap weights:



Data re-weighting.

Augment the problem with *data weights* $w_1, \dots, w_N$. We can write $\mathbb{E}_{\mathcal{P}(\theta|X;w)} [f(\theta)]$.


$$\ell_n(\theta) := \log \mathcal{P}(x_n|\theta)$$

$$\log \mathcal{P}(X|\theta; w) = \sum_{n=1}^N w_n \ell_n(\theta)$$



Original weights:



Leave-one-out weights:



Bootstrap weights:




Data re-weighting.

Augment the problem with *data weights* $w_1, \dots, w_N$. We can write $\mathbb{E}_{\mathcal{P}(\theta|X;w)} [f(\theta)]$.


$$\ell_n(\theta) := \log \mathcal{P}(x_n|\theta)$$

$$\log \mathcal{P}(X|\theta; w) = \sum_{n=1}^N w_n \ell_n(\theta)$$



Original weights:



Leave-one-out weights:



Bootstrap weights:



Data re-weighting.

Augment the problem with *data weights* $w_1, \dots, w_N$. We can write $\mathbb{E}_{\mathcal{P}(\theta|X;w)} [f(\theta)]$.


$$\ell_n(\theta) := \log \mathcal{P}(x_n|\theta)$$

$$\log \mathcal{P}(X|\theta; w) = \sum_{n=1}^N w_n \ell_n(\theta)$$



Original weights:



Leave-one-out weights:



Bootstrap weights:




Data re-weighting.

Augment the problem with *data weights* $w_1, \dots, w_N$. We can write $\mathbb{E}_{\mathcal{P}(\theta|X;w)} [f(\theta)]$.


$$\ell_n(\theta) := \log \mathcal{P}(x_n|\theta)$$

$$\log \mathcal{P}(X|\theta; w) = \sum_{n=1}^N w_n \ell_n(\theta)$$



Original weights:



Leave-one-out weights:



Bootstrap weights:




Data re-weighting.

Augment the problem with *data weights* $w_1, \dots, w_N$. We can write $\mathbb{E}_{\mathcal{P}(\theta|X;w)} [f(\theta)]$.


$$\ell_n(\theta) := \log \mathcal{P}(x_n|\theta)$$

$$\log \mathcal{P}(X|\theta; w) = \sum_{n=1}^N w_n \ell_n(\theta)$$



Original weights:



Leave-one-out weights:



Bootstrap weights:



The re-scaled slope $N\psi_n$ is known as the “influence function” at data point x_n .

$$\mathbb{E}_{\mathcal{P}(\theta|X;w)} [f(\theta)] - \mathbb{E}_{\mathcal{P}(\theta|X)} [f(\theta)] = \sum_{n=1}^N \psi_n (w_n - 1) + \mathcal{E}_n(w)$$

How can we use the approximation?

Assume the **slope** is computable and **error** is small.

$$\mathbb{E}_{p(\theta|X;w)} [f(\theta)] - \mathbb{E}_{p(\theta|X)} [f(\theta)] = \sum_{n=1}^N \psi_n(w_n - 1) + \mathcal{E}_n(w)$$

How can we use the approximation?

Assume the **slope** is computable and **error** is small.

$$\mathbb{E}_{p(\theta|X;w)} [f(\theta)] - \mathbb{E}_{p(\theta|X)} [f(\theta)] = \sum_{n=1}^N \psi_n(w_n - 1) + \mathcal{E}_n(w)$$

Bootstrap. Draw bootstrap weights $w \sim \mathcal{P}(w) = \text{Multinomial}(N, N^{-1})$.

$$\text{Bootstrap variance} = \text{Var}_{\mathcal{P}(w)} \left(\mathbb{E}_{p(\theta|X;w)} [f(\theta)] \right)$$

How can we use the approximation?

Assume the **slope** is computable and **error** is small.

$$\mathbb{E}_{p(\theta|X;w)} [f(\theta)] - \mathbb{E}_{p(\theta|X)} [f(\theta)] = \sum_{n=1}^N \psi_n(w_n - 1) + \mathcal{E}_n(w)$$

Bootstrap. Draw bootstrap weights $w \sim \mathcal{P}(w) = \text{Multinomial}(N, N^{-1})$.

$$\begin{aligned} \text{Bootstrap variance} &= \text{Var}_{\mathcal{P}(w)} \left(\mathbb{E}_{p(\theta|X;w)} [f(\theta)] \right) \\ &= \text{Var}_{\mathcal{P}(w)} \left(\sum_{n=1}^N \psi_n(w_n - 1) + \mathcal{E}_n(w) \right) \\ &= \frac{1}{N^2} \sum_{n=1}^N \left(\psi_n - \bar{\psi} \right)^2 + \text{Term involving } \mathcal{E}_n(w) \text{ for } n = 1, \dots, N \\ &\approx \frac{1}{N} \underbrace{\left(\frac{1}{N} \sum_{n=1}^N \left(\psi_n - \bar{\psi} \right)^2 \right)}_{\text{"Infinitesimal jackknife variance estimate"}} \end{aligned}$$

Bayesian von-Mises Expansion

Define the “generalized posterior” functional [Giordano and Broderick, 2026]:

$$T(\mathcal{P}, N) := \frac{\int g(\theta) \exp \left(N \int \ell(y|x, \theta) \mathcal{P}(x, y) dx dy \right) \pi(\theta) d\theta}{\int \exp \left(N \int \ell(y|x, \theta) \mathcal{P}(x, y) dx dy \right) \pi(\theta) d\theta}.$$

Recall that $\mathbb{E}_{\mathcal{P}(\theta|\text{Data})} [g(\theta)] = T(\hat{\mathcal{P}}_S, N_S)$

To derive frequentist variance estimates, study the *von Mises expansion*.

Let $\mathcal{P}_t = t\hat{\mathcal{P}}_S + (1 - t)\mathcal{P}_S$. Then form the series expansion in t :

$$\begin{aligned} \sqrt{N_S} \left(\mathbb{E}_{\mathcal{P}(\theta|\text{Data})} [g(\theta)] - T(\mathcal{P}_S, N_S) \right) &= \sqrt{N_S} \left(T(\hat{\mathcal{P}}_S, N_S) - T(\mathcal{P}_S, N_S) \right) \\ &= \sqrt{N_S} \left. \frac{\partial T(\mathcal{P}_t, N_S)}{\partial t} \right|_{t=0} (1 - 0) + \mathcal{E}(\tilde{t}) \\ &= \underbrace{\sqrt{N_S} \sum_{n=1}^N (\psi_i - \bar{\psi})}_{\text{Infinitesimal jackknife estimator}} + o_p(1) + \mathcal{E}(\tilde{t}). \end{aligned}$$

where

$$\psi_i := N_S \text{Cov}_{\mathcal{P}(\theta|\text{Data})} (g(\theta), \ell(y_i|x_i)) \quad \text{and} \quad \tilde{t} \in [0, 1].$$

Bayesian von-Mises Expansion Results

$$\sqrt{N_S} \left(\mathbb{E}_{\mathcal{P}(\theta|\text{Data})} [g(\theta)] - T(\mathcal{P}_S, N_S) \right) = \sqrt{N_S} \sum_{n=1}^N (\psi_i - \bar{\psi}) + o_p(1) + \mathcal{E}(\tilde{t}).$$

- Assume (sketch): A well-behaved MAP *maximum a posteriori* estimator $\hat{\theta}$ exists.
 - The dimension of θ is fixed as $N \rightarrow \infty$
 - The expected log likelihood has a strict maximum at θ_0
 - The observed log likelihood satisfies $\hat{\theta} \rightarrow \theta_0$
 - The expected log likelihood Hessian is negative definite at θ_0
- Assume (sketch): We can apply standard asymptotics.
 - The log prior and log likelihood are four times continuously differentiable
 - The prior is proper, and a technical set of prior expectations are finite
 - The log likelihood and its derivatives are dominated by a square-integrable envelope function for all θ in a neighborhood of θ_∞ .

Theorem 3 [Giordano and Broderick, 2026] (sketch):

$$\sup_{\tilde{t} \in [0,1]} |\mathcal{E}(\tilde{t})| \xrightarrow{P} 0 \quad \text{in the Bayesian von-Mises expansion.}$$

Corollary: Let $\frac{1}{N_S} \sum_{i=1}^{N_S} (\psi_i - \bar{\psi})^2 \xrightarrow{P} V$. By the CLT applied to ψ_i ,

$$\sqrt{N_S} \left(\mathbb{E}_{\mathcal{P}(\theta|\text{Data})} [g(\theta)] - T(\mathcal{P}_S, N_S) \right) \xrightarrow{d} \mathcal{N}(0, V).$$

Recall that our covariate shift diagnostic is given by

$$\Delta(w, \hat{\mathcal{P}}_S) = \text{Cov}_{\mathcal{P}(\theta|\text{Data})} \left(\frac{1}{N_T} \sum_{j=1}^{N_T} m(\theta^\top x_j), \sum_{i=1}^{N_S} w(x_i) \ell(y_i | x_i, \theta) \right) = \sum_{i=1}^{N_S} \psi_i w(x_i).$$

In other words, we are checking whether the influence function aligns with any of the regressors.




Figure 3: MrP Posterior Mean




Figure 3: MrP Posterior Variance




Figure 3: Covariate Shift Diagnostic

Local robustness with flows

Why flows for sensitivity analysis?

The IJ is built around linearizing posterior expectations:

$$\mathbb{E}_{\mathcal{P}^\epsilon(\theta)} [g(\theta)] \approx \mathbb{E}_{\mathcal{P}} [g(\theta)] + \left. \frac{\partial \mathbb{E}_{\mathcal{P}^\epsilon(\theta)} [g(\theta)]}{\partial \epsilon} \right|_{\epsilon=0} (\epsilon - 0),$$

where the derivative is estimated — with noise — from MCMC draws.

This has some shortcomings:

- Potential MCMC variance, not obvious how to smooth
- Linear extrapolation can lead to impossible values (e.g. negative variances)
- Complex quantities like posterior quantities are hard to linearize
- Bayesians like to work with draws, not analytic expectations

Idea: Instead of linearizing expectations, learn a flow from $\mathcal{P}$ to $\mathcal{P}^\epsilon$.

There's a large literature on learning flows between distant distributions. Our problem is easier:

- Only need “local” flows
- Our perturbations have a lot of structure (e.g. bootstraps)

Consider a family of distributions $\mathcal{P}^\epsilon(\theta) \propto \mathcal{P}(\theta) \exp(\epsilon h(\theta))$ and a test function $\phi(\theta)$. Then, by interchange of integration and differentiation,

$$\left. \frac{\partial \mathbb{E}_{\mathcal{P}^\epsilon(\theta)} [\phi(\theta)]}{\partial \epsilon} \right|_{\epsilon=0} = \text{Cov}_{\mathcal{P}(\theta)} (\phi(\theta), h(\theta)) .$$

We want a flow $\theta \mapsto f^\epsilon(\theta)$ such that, for small ϵ , $f^\epsilon(\theta) \sim \mathcal{P}^\epsilon(\theta)$ and $f^0(\theta) = \theta$.

Suppose there is a scalar-valued potential function $v(\theta)$ such that

$$\begin{aligned} f^\epsilon(\theta) &= \theta + \epsilon \nabla v(\theta) + O(\epsilon^2) \quad \text{as } \epsilon \rightarrow 0. \text{ Then:} \\ \mathbb{E}_{\mathcal{P}^\epsilon(\theta)} [\phi(\theta)] &= \mathbb{E}_{\mathcal{P}(\theta)} [\phi(f^\epsilon(\theta))] \Rightarrow \\ \left. \frac{\partial}{\partial \epsilon} \right|_{\epsilon=0} \mathbb{E}_{\mathcal{P}^\epsilon(\theta)} [\phi(\theta)] &= \left. \frac{\partial}{\partial \epsilon} \right|_{\epsilon=0} \mathbb{E}_{\mathcal{P}(\theta)} [\phi(\theta + \epsilon \nabla v(\theta))] \\ \parallel &\parallel \\ \mathbb{E}_{\mathcal{P}(\theta)} [\phi(\theta) h(\theta)] &= \mathbb{E}_{\mathcal{P}(\theta)} [\nabla \phi(\theta)^\top \nabla v(\theta)] \end{aligned}$$

- When does such a $v(\theta)$ exist?
- How can we find it in practice?

Consider a family of distributions $\mathcal{P}^\epsilon(\theta) \propto \mathcal{P}(\theta) \exp(\epsilon h(\theta))$.

Integration by parts (with compactly supported ϕ) gives:

$$\begin{aligned}\mathbb{E}_{\mathcal{P}(\theta)} [\phi(\theta) h(\theta)] &= \mathbb{E}_{\mathcal{P}(\theta)} [\nabla \phi(\theta)^\top \nabla v(\theta)] \\ &= \mathbb{E}_{\mathcal{P}(\theta)} [\phi(\theta) (-\nabla v(\theta)^\top \nabla \log \mathcal{P}(\theta) - \Delta v(\theta))] \\ &=: \mathbb{E}_{\mathcal{P}(\theta)} [\phi(\theta) \mathcal{L}_{\mathcal{P}}(v)(\theta)] \quad \text{for} \\ \mathcal{L}_{\mathcal{P}}(v) &:= -\nabla v(\theta)^\top \nabla \log \mathcal{P}(\theta) - \Delta v(\theta).\end{aligned}$$

This is a weak formulation of the linear differential equation

$$\mathcal{L}_{\mathcal{P}}(v) = h.$$

Note that $\mathcal{L}_{\mathcal{P}}$ is the drift operator, generator of Langevin diffusion, trace of the Stein operator.

Theorem (Schewengber): Assume that $\log \mathcal{P}(\theta) \in C^2(\mathcal{R}^d)$ and

$$\lim_{\theta \rightarrow \infty} \left(\frac{1}{2} \|\nabla \log \mathcal{P}(\theta)\|^2 - \Delta \log \mathcal{P}(\theta) \right) = +\infty.$$

Then $\mathcal{L}_{\mathcal{P}}^{-1}$ is bounded as a map from $\dot{H}^1(p) \mapsto \dot{H}^1(p)$.

Corollary: If $\log \mathcal{P}(\theta|\text{Data})$ is sufficiently regular, $\mathbb{E}_{\mathcal{P}(\theta|\text{Data})} [\|\nabla h(\theta)\|_2^2] < \infty$, and $|\mathbb{E}_{\mathcal{P}(\theta|\text{Data})} [h(\theta)]| < \infty$, then there exists a differentiable flow potential $v(\theta)$ that locally approximates $\mathcal{P}^\epsilon(\theta|\text{Data})$.

R. Giordano and T. Broderick. The Bayesian infinitesimal jackknife for variance. *arXiv preprint arXiv:2305.06466*, 2026.