

Next-generation Bayesian local robustness

Ryan Giordano (rgiordano@berkeley.edu, UC Berkeley)

Berkeley Neyman Seminar


September 2026

The basic problem


We have a survey population, for whom we observe:

- Covariates x (e.g. race, gender, zip code, age, education level)
- Responses y (e.g. A binary response to “do you support candidate Z”)

We want the average response in a target population, in which we observe only covariates.



Observe (x_i, y_i) for $i = 1, \dots, N_S$
(This is the “Data”, $\mathcal{D}$)



Observe x_j for $j = 1, \dots, N_T$

The basic problem

We have a survey population, for whom we observe:


- Covariates x (e.g. race, gender, zip code, age, education level)
- Responses y (e.g. A binary response to “do you support candidate Z”)

We want the average response in a target population, in which we observe only covariates.

The problem is that the populations may be very different, maybe leading to bias.



Observe (x_i, y_i) for $i = 1, \dots, N_S$
(This is the “Data”, $\mathcal{D}$)



Observe x_j for $j = 1, \dots, N_T$

The basic problem

We have a survey population, for whom we observe:


- Covariates x (e.g. race, gender, zip code, age, education level)
- Responses y (e.g. A binary response to “do you support candidate Z”)

We want the average response in a target population, in which we observe only covariates.

The problem is that the populations may be very different, maybe leading to bias.



Observe (x_i, y_i) for $i = 1, \dots, N_S$
(This is the “Data”, $\mathcal{D}$)



Observe x_j for $j = 1, \dots, N_T$

This talk will study “Multilevel regression and poststratification” estimators:

- Set $\mathbb{E}[y|x, \theta] = m(\theta^\top x)$ and prior $\mathcal{P}(\theta)$ (e.g. Hierarchical logistic regression)
- Estimate the posterior $\mathcal{P}(\theta|\mathcal{D})$ with Markov Chain Monte Carlo (MCMC)
- Take $\hat{\mu} = \frac{1}{N_T} \sum_{j=1}^{N_T} \hat{y}_j$ where $\hat{y}_j = \mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})}[y|x_j]$.

The basic problem

We have a survey population, for whom we observe:


- Covariates x (e.g. race, gender, zip code, age, education level)
- Responses y (e.g. A binary response to “do you support candidate Z”)

We want the average response in a target population, in which we observe only covariates.

The problem is that the populations may be very different, maybe leading to bias.



Observe (x_i, y_i) for $i = 1, \dots, N_S$
(This is the “Data”, $\mathcal{D}$)



Observe x_j for $j = 1, \dots, N_T$

This talk will study “Multilevel regression and poststratification” estimators:

- Set $\mathbb{E}[y|x, \theta] = m(\theta^\top x)$ and prior $\mathcal{P}(\theta)$ (e.g. Hierarchical logistic regression)
- Estimate the posterior $\mathcal{P}(\theta|\mathcal{D})$ with Markov Chain Monte Carlo (MCMC)
- Take $\hat{\mu} = \frac{1}{N_T} \sum_{j=1}^{N_T} \hat{y}_j$ where $\hat{y}_j = \mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})}[y|x_j]$.

Problem: MCMC is time consuming and the regression model is surely misspecified.

Goal: Meaningful regression specification checks with frequentist uncertainty.

Result preview




Figure 1: A covariate shift diagnostic

Here is our check for two models on the same simulated data: one correctly specified and one misspecified. The misspecified model is more extreme.

Problem: What is “large enough” to be real evidence of misspecification?

Possible answer: Frequentist confidence intervals via the parametric bootstrap.

Result preview




Figure 1: A covariate shift diagnostic

The frequentist interval based on the parametric bootstrap rejects the misspecified model but not the correctly specified model.

Problem: Expensive: each bootstrap sample required a new MCMC run.

Possible answer: Infinitesimal jackknife confidence intervals? [Giordano and Broderick, 2026]

Result preview




Figure 1: A covariate shift diagnostic

IJ confidence intervals are fast and often work well for posterior means.
But they are not accurate for this more complex diagnostic (high Monte Carlo error).

Result preview




Figure 1: A covariate shift diagnostic

IJ confidence intervals are fast and often work well for posterior means. But they are not accurate for this more complex diagnostic (high Monte Carlo error).

Answer: Our flow-based local sensitivity measures (“Otto grad”) produce good approximations to the bootstrap at negligible computational expense.

- Bayesian local regression specification checks
 - (With Alice Cima, Erin Hartman, Jared Miller, Avi Feller)
- Frequentist variability of Bayesian posteriors with the infinitesimal jackknife
 - (With Tamara Broderick)
- Flow-based Bayesian local robustness: “Otto grad”
 - (With Lucas Schwengber)

TODO: put in pictures

Bayesian local regression specification checks

$$\hat{\mu} = \frac{1}{N_T} \sum_{j=1}^{N_T} \hat{y}_j = \underbrace{\mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})} \left[\frac{1}{N_T} \sum_{j=1}^{N_T} m(\theta^\top x_j) \right]}_{\text{Want to check specification without re-running MCMC}}.$$

Some existing local approaches:

- Specify an alternative model + test with Bayes factor [Cabral et al.]
 - Bayes factors can depend strongly on priors (Lindley's paradox)
 - Articulating a family of models is inconvenient
 - Doesn't take the target into account
- Use out-of-sample checking [Vehtari et al., Kuh et al.]
 - The task is not to predict on a new IID survey datapoint

$$\hat{\mu} = \frac{1}{N_T} \sum_{j=1}^{N_T} \hat{y}_j = \underbrace{\mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})} \left[\frac{1}{N_T} \sum_{j=1}^{N_T} m(\theta^\top x_j) \right]}_{\text{Want to check specification without re-running MCMC}}.$$

Some existing local approaches:

- Specify an alternative model + test with Bayes factor [Cabral et al.]
 - Bayes factors can depend strongly on priors (Lindley's paradox)
 - Articulating a family of models is inconvenient
 - Doesn't take the target into account
- Use out-of-sample checking [Vehtari et al., Kuh et al.]
 - The task is not to predict on a new IID survey datapoint

The model is definitely misspecified, we want to check whether it's *importantly* misspecified.

Local specification checks

$$\hat{\mu} = \frac{1}{N_T} \sum_{j=1}^{N_T} \hat{y}_j = \underbrace{\mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})} \left[\frac{1}{N_T} \sum_{j=1}^{N_T} m(\theta^\top x_j) \right]}_{\substack{\text{Want to check specification} \\ \text{without re-running MCMC}}}.$$

Why do we care about correct specification? We want to use $\mathcal{P}_S(x)$ to learn about $\mathcal{P}_T(x)$.

Correct specification $\Rightarrow$ Regression estimates are invariant to $\mathcal{P}_S(x)$

Local specification checks

$$\hat{\mu} = \frac{1}{N_T} \sum_{j=1}^{N_T} \hat{y}_j = \underbrace{\mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})} \left[\frac{1}{N_T} \sum_{j=1}^{N_T} m(\theta^\top x_j) \right]}_{\substack{\text{Want to check specification} \\ \text{without re-running MCMC}}}.$$

Why do we care about correct specification? We want to use $\mathcal{P}_S(x)$ to learn about $\mathcal{P}_T(x)$.

Correct specification $\Rightarrow$ Regression estimates are invariant to $\mathcal{P}_S(x)$
This is what we actually want — let's check this directly.

$$\hat{\mu} = \frac{1}{N_T} \sum_{j=1}^{N_T} \hat{y}_j = \underbrace{\mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})} \left[\frac{1}{N_T} \sum_{j=1}^{N_T} m(\theta^\top x_j) \right]}_{\substack{\text{Want to check specification} \\ \text{without re-running MCMC}}}.$$

Why do we care about correct specification? We want to use $\mathcal{P}_S(x)$ to learn about $\mathcal{P}_T(x)$.

$$\text{Correct specification} \quad \Rightarrow \quad \underbrace{\text{Regression estimates are invariant to } \mathcal{P}_S(x)}_{\text{This is what we actually want — let's check this directly.}}$$

Let $\mathcal{P}(x, y)$ denote a generic joint density of x, y .

Let $\hat{\mathcal{P}}_S$ the empirical distribution on the survey data.

1. Write our estimand as a functional $\mu(\mathcal{P})$ that takes in a density and returns a MrP estimate
2. Define a covariate density perturbation $w(x)$
3. Define $\mathcal{P}_\epsilon(x, y) = \mathcal{P}(x, y) + \epsilon w(x) \mathcal{P}(y|x) = (\mathcal{P}(x) + \epsilon w(x)) \mathcal{P}(y|x)$
4. Compute the derivative $\Delta(w, \mathcal{P}) = \partial \mu(\mathcal{P}_\epsilon) / \partial \epsilon$
5. Plug the empirical distribution into $\Delta(w, \hat{\mathcal{P}}_S)$ and check whether ≈ 0

Local specification checks

$$\hat{\mu} = \frac{1}{N_T} \sum_{j=1}^{N_T} \hat{y}_j = \underbrace{\mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})} \left[\frac{1}{N_T} \sum_{j=1}^{N_T} m(\theta^\top x_j) \right]}_{\text{Want to check specification without re-running MCMC}}.$$

Why do we care about correct specification? We want to use $\mathcal{P}_S(x)$ to learn about $\mathcal{P}_T(x)$.

Correct specification $\Rightarrow$ Regression estimates are invariant to $\mathcal{P}_S(x)$
This is what we actually want — let's check this directly.

Let $\mathcal{P}(x, y)$ denote a generic joint density of x, y .

Let $\hat{\mathcal{P}}_S$ the empirical distribution on the survey data.

1. Write our estimand as a functional $\mu(\mathcal{P})$ that takes in a density and returns a MrP estimate
2. Define a covariate density perturbation $w(x)$
3. Define $\mathcal{P}_\epsilon(x, y) = \mathcal{P}(x, y) + \epsilon w(x)\mathcal{P}(y|x) = (\mathcal{P}(x) + \epsilon w(x))\mathcal{P}(y|x)$
4. Compute the derivative $\Delta(w, \mathcal{P}) = \partial\mu(\mathcal{P}_\epsilon)/\partial\epsilon$
5. Plug the empirical distribution into $\Delta(w, \hat{\mathcal{P}}_S)$ and check whether ≈ 0

Lots of similar ideas in causal inference, kernel learning, transfer learning [Bühlmann, Rojas-Carulla et al., Muandet et al.]. This is just a local Bayesian diagnostics version.

Simple visual example




Figure 2: A covariate shift diagnostic

Consider logistic regression with two regressors, x_1 and x_2 .

Simple visual example




Figure 2: A covariate shift diagnostic

Consider logistic regression with two regressors, x_1 and x_2 .

Taking $w(x) = \mathbb{I}(x_1 > 1)$, we upweight positive examples.

Simple visual example




Figure 2: A covariate shift diagnostic

Consider logistic regression with two regressors, x_1 and x_2 .

Taking $w(x) = \mathbb{I}(x_1 > 1)$, we upweight positive examples.

With enough data, a regression that omits x_1 will change its coefficient of x_2 under re-weighting, whereas a correctly specified regression will not need to.

Bayesian sensitivity analysis

Let the regression log likelihood be $\ell(y|x, \theta)$, and fix $w(x)$. Then define

$$\mathcal{P}(\theta|\mathcal{D}) \propto \exp\left(\sum_{i=1}^{N_S} \ell(y_i|x_i, \theta)\right) \pi(\theta) \text{ and } \mathcal{P}^\epsilon(\theta) \propto \mathcal{P}(\theta|\mathcal{D}) \exp\left(\sum_{i=1}^{N_S} \epsilon w(x_i) \ell(y_i|x_i, \theta)\right)$$

Bayesian sensitivity analysis

Let the regression log likelihood be $\ell(y|x, \theta)$, and fix $w(x)$. Then define

$$\mathcal{P}(\theta|\mathcal{D}) \propto \exp\left(\sum_{i=1}^{N_S} \ell(y_i|x_i, \theta)\right) \pi(\theta) \text{ and } \mathcal{P}^\epsilon(\theta) \propto \mathcal{P}(\theta|\mathcal{D}) \exp\left(\sum_{i=1}^{N_S} \epsilon w(x_i) \ell(y_i|x_i, \theta)\right)$$

Theorem (Giordano et al. [2018], others): Consider a family of distributions $\mathcal{P}^\epsilon(\theta) \propto \mathcal{P}(\theta) \exp(\epsilon h(\theta))$ and a function $\phi(\theta)$. By interchange of integration and differentiation,

$$\left. \frac{\partial \mathbb{E}_{\mathcal{P}^\epsilon(\theta)} [\phi(\theta)]}{\partial \epsilon} \right|_{\epsilon=0} = \text{Cov}_{\mathcal{P}(\theta)} (\phi(\theta), h(\theta)) .$$

Bayesian sensitivity analysis

Let the regression log likelihood be $\ell(y|x, \theta)$, and fix $w(x)$. Then define

$$\mathcal{P}(\theta|\mathcal{D}) \propto \exp\left(\sum_{i=1}^{N_S} \ell(y_i|x_i, \theta)\right) \pi(\theta) \text{ and } \mathcal{P}^\epsilon(\theta) \propto \mathcal{P}(\theta|\mathcal{D}) \exp\left(\sum_{i=1}^{N_S} \epsilon w(x_i) \ell(y_i|x_i, \theta)\right)$$

Theorem (Giordano et al. [2018], others): Consider a family of distributions $\mathcal{P}^\epsilon(\theta) \propto \mathcal{P}(\theta) \exp(\epsilon h(\theta))$ and a function $\phi(\theta)$. By interchange of integration and differentiation,

$$\left. \frac{\partial \mathbb{E}_{\mathcal{P}^\epsilon(\theta)} [\phi(\theta)]}{\partial \epsilon} \right|_{\epsilon=0} = \text{Cov}_{\mathcal{P}(\theta)} (\phi(\theta), h(\theta)) .$$

Our diagnostic is then given by

$$\Delta(w, \mathcal{D}) := \left. \frac{\partial \hat{\mu}}{\partial \epsilon} \right|_{\epsilon=0} = \text{Cov}_{\mathcal{P}(\theta|\mathcal{D})} \left(\frac{1}{N_T} \sum_{j=1}^{N_T} m(\theta^\top x_j), \sum_{i=1}^{N_S} w(x_i) \ell(y_i|x_i, \theta) \right)$$

We expect $\Delta(w, \mathcal{D})$ to be *close* to zero when the model is correctly specified.

Bayesian sensitivity analysis

Let the regression log likelihood be $\ell(y|x, \theta)$, and fix $w(x)$. Then define

$$\mathcal{P}(\theta|\mathcal{D}) \propto \exp \left(\sum_{i=1}^{N_S} \ell(y_i|x_i, \theta) \right) \pi(\theta) \text{ and } \mathcal{P}^\epsilon(\theta) \propto \mathcal{P}(\theta|\mathcal{D}) \exp \left(\sum_{i=1}^{N_S} \epsilon w(x_i) \ell(y_i|x_i, \theta) \right)$$

Theorem (Giordano et al. [2018], others): Consider a family of distributions $\mathcal{P}^\epsilon(\theta) \propto \mathcal{P}(\theta) \exp(\epsilon h(\theta))$ and a function $\phi(\theta)$. By interchange of integration and differentiation,

$$\left. \frac{\partial \mathbb{E}_{\mathcal{P}^\epsilon(\theta)} [\phi(\theta)]}{\partial \epsilon} \right|_{\epsilon=0} = \text{Cov}_{\mathcal{P}(\theta)} (\phi(\theta), h(\theta)) .$$

Our diagnostic is then given by

$$\Delta(w, \mathcal{D}) := \left. \frac{\partial \hat{\mu}}{\partial \epsilon} \right|_{\epsilon=0} = \text{Cov}_{\mathcal{P}(\theta|\mathcal{D})} \left(\frac{1}{N_T} \sum_{j=1}^{N_T} m(\theta^\top x_j), \sum_{i=1}^{N_S} w(x_i) \ell(y_i|x_i, \theta) \right)$$

We expect $\Delta(w, \mathcal{D})$ to be *close* to zero when the model is correctly specified.

To assess closeness to zero, compare with $\Delta(w, \mathcal{D}^*)$, where $\mathcal{D}^*$ is a parametric bootstrap.




Figure 3: A covariate shift diagnostic

Simulation data: Logistic regression with five active regressors, each with covariate shift.

Model 0 (correct) : $y \sim \text{Logistic}(1 + X_1 + X_2 + X_3 + X_4 + X_5)$

Model 1 (misspecified) : $y \sim \text{Logistic}(1 + X_3 + X_4 + X_5)$

Compute diagnostic with $w(x) = \mathbb{I}(X_2 > 0.878)$

Frequentist variance with the infinitesimal
jackknife

Why frequentist variance?

For the remainder of the talk, focus on approximating the *frequentist* sampling distribution under $x_i, y_i \stackrel{iid}{\sim} \mathcal{P}_S(x, y)$ of quantities like our diagnostic:

$$\Delta = \text{Cov}_{\mathcal{P}(\theta|\mathcal{D})} \left(\frac{1}{N_T} \sum_{j=1}^{N_T} m(\theta^\top x_j), \sum_{i=1}^{N_S} w(x_i) \ell(y_i | x_i, \theta) \right)$$

No need to commit to our covariate shift diagnostic to care!

- Frequentist variability of *any* model checking diagnostic is interesting
- Frequentist variability remains meaningful under model misspecification (Huggins and Miller, IJ papers)

We will form *local approximations* computable only from *the original posterior*.

We will compare with computationally expensive bootstrap:

- Compute Δ on the original dataset
- For $B = 1, \dots, 100$:
 1. Sample $(x_1^*, y_1^*), \dots, (x_{N_S}^*, y_{N_S}^*)$ (parametric or nonparametric bootstrap)
 2. Run MCMC to estimate $\mathcal{P}(\theta | (x_1^*, y_1^*), \dots, (x_{N_S}^*, y_{N_S}^*))$
 3. Use the draws to estimate Δ_B^*
- Compare Δ with quantiles of $\Delta_B^* - \Delta$

Data re-weighting.

Augment the problem with *data weights* $w_1, \dots, w_N$. We can write $\mathbb{E}_{\mathcal{P}(\theta|X;w)} [f(\theta)]$.

$$\ell_n(\theta) := \log \mathcal{P}(x_n|\theta) \qquad \log \mathcal{P}(X|\theta; w) = \sum_{n=1}^N w_n \ell_n(\theta)$$

Data re-weighting.

Augment the problem with *data weights* $w_1, \dots, w_N$. We can write $\mathbb{E}_{\mathcal{P}(\theta|X;w)} [f(\theta)]$.

$$\ell_n(\theta) := \log \mathcal{P}(x_n|\theta) \qquad \log \mathcal{P}(X|\theta; w) = \sum_{n=1}^N w_n \ell_n(\theta)$$

Original weights:




Data re-weighting.


Augment the problem with *data weights* $w_1, \dots, w_N$. We can write $\mathbb{E}_{\mathcal{P}(\theta|X;w)} [f(\theta)]$.

$$\ell_n(\theta) := \log \mathcal{P}(x_n|\theta) \qquad \log \mathcal{P}(X|\theta; w) = \sum_{n=1}^N w_n \ell_n(\theta)$$

Original weights:



Leave-one-out weights:



Data re-weighting.

Augment the problem with *data weights* $w_1, \dots, w_N$. We can write $\mathbb{E}_{\mathcal{P}(\theta|X;w)} [f(\theta)]$.

$$\ell_n(\theta) := \log \mathcal{P}(x_n|\theta)$$

$$\log \mathcal{P}(X|\theta; w) = \sum_{n=1}^N w_n \ell_n(\theta)$$


Original weights:



Leave-one-out weights:



Bootstrap weights:




Data re-weighting.

Augment the problem with *data weights* $w_1, \dots, w_N$. We can write $\mathbb{E}_{\mathcal{P}(\theta|X;w)} [f(\theta)]$.


$$\ell_n(\theta) := \log \mathcal{P}(x_n|\theta)$$

$$\log \mathcal{P}(X|\theta; w) = \sum_{n=1}^N w_n \ell_n(\theta)$$



Original weights:



Leave-one-out weights:



Bootstrap weights:



Data re-weighting.

Augment the problem with *data weights* $w_1, \dots, w_N$. We can write $\mathbb{E}_{\mathcal{P}(\theta|X;w)} [f(\theta)]$.


$$\ell_n(\theta) := \log \mathcal{P}(x_n|\theta)$$

$$\log \mathcal{P}(X|\theta; w) = \sum_{n=1}^N w_n \ell_n(\theta)$$



Original weights:



Leave-one-out weights:



Bootstrap weights:




Data re-weighting.

Augment the problem with *data weights* $w_1, \dots, w_N$. We can write $\mathbb{E}_{\mathcal{P}(\theta|X;w)} [f(\theta)]$.


$$\ell_n(\theta) := \log \mathcal{P}(x_n|\theta)$$

$$\log \mathcal{P}(X|\theta; w) = \sum_{n=1}^N w_n \ell_n(\theta)$$



Original weights:



Leave-one-out weights:



Bootstrap weights:




Data re-weighting.

Augment the problem with *data weights* $w_1, \dots, w_N$. We can write $\mathbb{E}_{\mathcal{P}(\theta|X;w)} [f(\theta)]$.


$$\ell_n(\theta) := \log \mathcal{P}(x_n|\theta)$$

$$\log \mathcal{P}(X|\theta; w) = \sum_{n=1}^N w_n \ell_n(\theta)$$



Original weights:



Leave-one-out weights:



Bootstrap weights:




Data re-weighting.

Augment the problem with *data weights* $w_1, \dots, w_N$. We can write $\mathbb{E}_{\mathcal{P}(\theta|X;w)} [f(\theta)]$.


$$\ell_n(\theta) := \log \mathcal{P}(x_n|\theta)$$

$$\log \mathcal{P}(X|\theta; w) = \sum_{n=1}^N w_n \ell_n(\theta)$$



Original weights:



Leave-one-out weights:



Bootstrap weights:




Data re-weighting.

Augment the problem with *data weights* $w_1, \dots, w_N$. We can write $\mathbb{E}_{\mathcal{P}(\theta|X;w)} [f(\theta)]$.


$$\ell_n(\theta) := \log \mathcal{P}(x_n|\theta)$$

$$\log \mathcal{P}(X|\theta; w) = \sum_{n=1}^N w_n \ell_n(\theta)$$



Original weights:



Leave-one-out weights:



Bootstrap weights:




Data re-weighting.

Augment the problem with *data weights* $w_1, \dots, w_N$. We can write $\mathbb{E}_{\mathcal{P}(\theta|X;w)} [f(\theta)]$.


$$\ell_n(\theta) := \log \mathcal{P}(x_n|\theta)$$

$$\log \mathcal{P}(X|\theta; w) = \sum_{n=1}^N w_n \ell_n(\theta)$$



Original weights:



Leave-one-out weights:



Bootstrap weights:



The re-scaled slope $N\psi_n$ is known as the “influence function” at data point x_n .

$$\mathbb{E}_{p(\theta|X;w)} [f(\theta)] - \mathbb{E}_{p(\theta|X)} [f(\theta)] = \sum_{n=1}^N \psi_n (w_n - 1) + \mathcal{E}_n(w)$$

How can we use the approximation?

Assume the **slope** is computable and **error** is small.

$$\mathbb{E}_{p(\theta|X;w)} [f(\theta)] - \mathbb{E}_{p(\theta|X)} [f(\theta)] = \sum_{n=1}^N \psi_n(w_n - 1) + \mathcal{E}_n(w)$$

How can we use the approximation?

Assume the **slope** is computable and **error** is small.

$$\mathbb{E}_{p(\theta|X;w)} [f(\theta)] - \mathbb{E}_{p(\theta|X)} [f(\theta)] = \sum_{n=1}^N \psi_n(w_n - 1) + \mathcal{E}_n(w)$$

Bootstrap. Draw bootstrap weights $w \sim \mathcal{P}(w) = \text{Multinomial}(N, N^{-1})$.

$$\text{Bootstrap variance} = \text{Var}_{\mathcal{P}(w)} \left(\mathbb{E}_{p(\theta|X;w)} [f(\theta)] \right)$$

How can we use the approximation?

Assume the **slope** is computable and **error** is small.

$$\mathbb{E}_{p(\theta|X;w)} [f(\theta)] - \mathbb{E}_{p(\theta|X)} [f(\theta)] = \sum_{n=1}^N \psi_n(w_n - 1) + \mathcal{E}_n(w)$$

Bootstrap. Draw bootstrap weights $w \sim \mathcal{P}(w) = \text{Multinomial}(N, N^{-1})$.

$$\begin{aligned} \text{Bootstrap variance} &= \text{Var}_{\mathcal{P}(w)} \left(\mathbb{E}_{p(\theta|X;w)} [f(\theta)] \right) \\ &= \text{Var}_{\mathcal{P}(w)} \left(\sum_{n=1}^N \psi_n(w_n - 1) + \mathcal{E}_n(w) \right) \\ &= \frac{1}{N^2} \sum_{n=1}^N \left(\psi_n - \bar{\psi} \right)^2 + \text{Term involving } \mathcal{E}_n(w) \text{ for } n = 1, \dots, N \\ &\approx \frac{1}{N} \underbrace{\left(\frac{1}{N} \sum_{n=1}^N \left(\psi_n - \bar{\psi} \right)^2 \right)}_{\text{"Infinitesimal jackknife variance estimate"}} \end{aligned}$$

Bayesian von-Mises Expansion

Define the “generalized posterior” functional [Giordano and Broderick, 2026]:

$$T(\mathcal{P}, N) := \frac{\int g(\theta) \exp \left(N \int \ell(y|x, \theta) \mathcal{P}(x, y) dx dy \right) \pi(\theta) d\theta}{\int \exp \left(N \int \ell(y|x, \theta) \mathcal{P}(x, y) dx dy \right) \pi(\theta) d\theta}.$$

Recall that $\mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})} [g(\theta)] = T(\hat{\mathcal{P}}_S, N_S)$

To derive frequentist variance estimates, study the *von Mises expansion*.

Let $\mathcal{P}_t = t\hat{\mathcal{P}}_S + (1 - t)\mathcal{P}_S$. Then form the series expansion in t :

$$\begin{aligned} \sqrt{N_S} \left(\mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})} [g(\theta)] - T(\mathcal{P}_S, N_S) \right) &= \sqrt{N_S} \left(T(\hat{\mathcal{P}}_S, N_S) - T(\mathcal{P}_S, N_S) \right) \\ &= \sqrt{N_S} \left. \frac{\partial T(\mathcal{P}_t, N_S)}{\partial t} \right|_{t=0} (1 - 0) + \mathcal{E}(\tilde{t}) \\ &= \underbrace{\sqrt{N_S} \sum_{n=1}^N (\psi_i - \bar{\psi})}_{\text{Infinitesimal jackknife estimator}} + o_p(1) + \mathcal{E}(\tilde{t}). \end{aligned}$$

where

$$\psi_i := N_S \text{Cov}_{\mathcal{P}(\theta|\mathcal{D})} (g(\theta), \ell(y_i|x_i)) \quad \text{and} \quad \tilde{t} \in [0, 1].$$

Bayesian von-Mises Expansion Results

$$\sqrt{N_S} \left(\mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})} [g(\theta)] - T(\mathcal{P}_S, N_S) \right) = \sqrt{N_S} \sum_{n=1}^N (\psi_i - \bar{\psi}) + o_p(1) + \mathcal{E}(\tilde{t}).$$

- Assume (sketch): A well-behaved MAP *maximum a posteriori* estimator $\hat{\theta}$ exists.
 - The dimension of θ is fixed as $N \rightarrow \infty$
 - The expected log likelihood has a strict maximum at θ_0
 - The observed log likelihood satisfies $\hat{\theta} \rightarrow \theta_0$
 - The expected log likelihood Hessian is negative definite at θ_0
- Assume (sketch): We can apply standard asymptotics.
 - The log prior and log likelihood are four times continuously differentiable
 - The prior is proper, and a technical set of prior expectations are finite
 - The log likelihood and its derivatives are dominated by a square-integrable envelope function for all θ in a neighborhood of θ_∞ .

Theorem 3 [Giordano and Broderick, 2026] (sketch):

$$\sup_{\tilde{t} \in [0,1]} |\mathcal{E}(\tilde{t})| \xrightarrow{P} 0 \quad \text{in the Bayesian von-Mises expansion.}$$

Corollary: Let $\frac{1}{N_S} \sum_{i=1}^{N_S} (\psi_i - \bar{\psi})^2 \xrightarrow{P} V$. By the CLT applied to ψ_i ,

$$\sqrt{N_S} \left(\mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})} [g(\theta)] - T(\mathcal{P}_S, N_S) \right) \xrightarrow{d} \mathcal{N}(0, V).$$

Recall that our covariate shift diagnostic is given by

$$\Delta(w, \hat{\mathcal{P}}_S) = \text{Cov}_{\mathcal{P}(\theta|\mathcal{D})} \left(\frac{1}{N_T} \sum_{j=1}^{N_T} m(\theta^\top x_j), \sum_{i=1}^{N_S} w(x_i) \ell(y_i | x_i, \theta) \right) = \sum_{i=1}^{N_S} \psi_i w(x_i).$$

In other words, we are checking whether the influence function aligns with any of the regressors.

To compute the parametric bootstrap, it's a simple change. (Cite MrP as a special case.)




Figure 4: MrP Posterior Mean




Figure 4: MrP Posterior Variance

Data re-weighting.




Figure 4: Covariate Shift Diagnostic

Need something better for noisier / more complex posterior quantities!

Local robustness with flows

Why flows for sensitivity analysis?

The IJ is built around linearizing posterior expectations:

$$\mathbb{E}_{\mathcal{P}^\epsilon(\theta)} [g(\theta)] \approx \mathbb{E}_{\mathcal{P}} [g(\theta)] + \left. \frac{\partial \mathbb{E}_{\mathcal{P}^\epsilon(\theta)} [g(\theta)]}{\partial \epsilon} \right|_{\epsilon=0} (\epsilon - 0),$$

where the derivative is estimated — with noise — from MCMC draws.

This has some shortcomings:

- Potential MCMC variance, not obvious how to smooth
- Linear extrapolation can lead to impossible values (e.g. negative variances)
- Complex quantities like posterior quantiles are hard to linearize
- Bayesians like to work with draws, not analytic expectations (CITE)

Idea: Instead of linearizing expectations, learn a flow from $\mathcal{P}$ to $\mathcal{P}^\epsilon$.

There's a large literature on learning flows between distant distributions. (CITE)

Our problem is easier:

- Only need “local” flows
- Our perturbations have a lot of structure (e.g. bootstraps)

Consider a family of distributions $\mathcal{P}^\epsilon(\theta) \propto \mathcal{P}(\theta) \exp(\epsilon h(\theta))$ and a test function $\phi(\theta)$. Then, by interchange of integration and differentiation,

$$\left. \frac{\partial \mathbb{E}_{\mathcal{P}^\epsilon(\theta)} [\phi(\theta)]}{\partial \epsilon} \right|_{\epsilon=0} = \text{Cov}_{\mathcal{P}(\theta)} (\phi(\theta), h(\theta)) .$$

We want a flow $\theta \mapsto f^\epsilon(\theta)$ such that, for small ϵ , $f^\epsilon(\theta) \sim \mathcal{P}^\epsilon(\theta)$ and $f^0(\theta) = \theta$.

Suppose there is a scalar-valued potential function $v(\theta)$ such that

$$\begin{aligned} f^\epsilon(\theta) &= \theta + \epsilon \nabla v(\theta) + O(\epsilon^2) \quad \text{as } \epsilon \rightarrow 0. \text{ Then:} \\ \mathbb{E}_{\mathcal{P}^\epsilon(\theta)} [\phi(\theta)] &= \mathbb{E}_{\mathcal{P}(\theta)} [\phi(f^\epsilon(\theta))] \Rightarrow \\ \left. \frac{\partial}{\partial \epsilon} \right|_{\epsilon=0} \mathbb{E}_{\mathcal{P}^\epsilon(\theta)} [\phi(\theta)] &= \left. \frac{\partial}{\partial \epsilon} \right|_{\epsilon=0} \mathbb{E}_{\mathcal{P}(\theta)} [\phi(\theta + \epsilon \nabla v(\theta))] \\ \parallel &\parallel \\ \mathbb{E}_{\mathcal{P}(\theta)} [\phi(\theta) h(\theta)] &= \mathbb{E}_{\mathcal{P}(\theta)} [\nabla \phi(\theta)^\top \nabla v(\theta)] \end{aligned}$$

- When does such a $v(\theta)$ exist?
- How can we find it in practice?

When does such a $v(\theta)$ exist?

Consider a family of distributions $\mathcal{P}^\epsilon(\theta) \propto \mathcal{P}(\theta) \exp(\epsilon h(\theta))$ and a test function $\phi(\theta)$. Then, by interchange of integration and differentiation,

$$\left. \frac{\partial \mathbb{E}_{\mathcal{P}^\epsilon(\theta)} [\phi(\theta)]}{\partial \epsilon} \right|_{\epsilon=0} = \text{Cov}_{\mathcal{P}(\theta)} (\phi(\theta), h(\theta)) .$$

We want a flow $\theta \mapsto f^\epsilon(\theta)$ such that, for small ϵ , $f^\epsilon(\theta) \sim \mathcal{P}^\epsilon(\theta)$ and $f^0(\theta) = \theta$.

When does such a $v(\theta)$ exist?

Theorem (Schewengber 2026): Assume that

- $\log \mathcal{P}(\theta) \in C^2(\mathcal{R}^d)$
- $\lim_{\theta \rightarrow \infty} \left(\frac{1}{2} \|\nabla \log \mathcal{P}(\theta)\|^2 - \Delta \log \mathcal{P}(\theta) \right) = +\infty$.
- $\mathbb{E}_{\mathcal{P}(\theta)} [h(\theta)^2] < \infty$

Then there exists a solution v to the differential equation

$$-\mathcal{L}_{\mathcal{P}}(v) := \nabla v^\top \nabla \log \mathcal{P} + \Delta v = h$$

such that $\|\nabla v\|_{L^2(\mathcal{P})} < \infty$, and the flow $\theta_\epsilon = \theta + \epsilon \nabla v(\theta)$ satisfies $\theta_\epsilon \sim \mathcal{P}^\epsilon(\theta)$ for small ϵ .

Note that $\mathcal{L}_{\mathcal{P}}$ is the drift operator, generator of Langevin diffusion, trace of the Stein operator.

How can we find $v(\theta)$ in practice?

Fix $h(\theta)$. Specify

- A family of real-valued test functions $\phi_n(\theta)$ for $n \in [N]$
- A family of real-valued basis functions $b_m(\theta)$ for $m \in [M]$

Let $v(\theta; \beta) := \sum_{m=1}^M \beta_m b_m(\theta)$. Then

$$\underbrace{\mathbb{E}_{\mathcal{P}(\theta)} [\phi_n(\theta) h(\theta)]}_{\text{Estimate with MCMC!}} = \sum_{m=1}^M \beta_m \underbrace{\mathbb{E}_{\mathcal{P}(\theta)} [\nabla \phi_n(\theta)^\top \nabla b_m(\theta)]}_{\text{Estimate with MCMC!}}$$

This gives N linear equations in M unknowns. Known as “Otto’s calculus” in differential equations. Formally similar to numerical PDE solutions known as Galerkin methods, except we have a stochastic grid ($\theta \sim \mathcal{P}(\theta)$) rather than a deterministic one.

Unlike standard linear regression:

- We are free to choose ϕ_n and b_m
- Both F and B are measured with noise
- We have a lot of information about the covariance of F and B

Theorem (Schewengber 2026, sketch): The solution gives the best L2 approximation to $\|\nabla v(\theta)\|_{L^2(\mathcal{P})}$ in the span of $\phi_n(\theta)$. When $\mathcal{P}$ is Gaussian, the ordered Hermite polynomials give the best finite-rank approximation.

Suppose we believe that $\mathcal{P}(\theta|\mathcal{D})$ is approximately Gaussian.

Take the basis functions $b_m(\theta) = \theta_m$ (1st Hermite polynomial) and test functions $\phi_n(\theta) = b_m(\theta)$. Then the optimal flow is given by

$$\theta_\epsilon = \theta + \epsilon \underbrace{\text{Cov}_{\mathcal{P}(\theta|\mathcal{D})}(\theta, h(\theta))}_{\left. \frac{\partial \mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})}[\theta]}{\partial \epsilon} \right|_{\epsilon=0}}$$

This just translates the posterior in the direction of the mean changes.

Suppose we have $\theta_m \sim \mathcal{P}(\theta|\mathcal{D})$, and bootstrap weights w_i . Set

$$\theta_m^* = \theta_m + \frac{1}{N_S} \sum_{i=1}^{N_S} w_i \psi_i$$

and compute any posterior quantities from these samples.




Figure 5: MrP Posterior Mean




Figure 5: MrP Posterior Variance




Figure 5: Covariate Shift Diagnostic

Understanding the gap between IJ and Otto is ongoing work!

- Bayesian local regression specification checks
 - Use invariance to covariate shift as a diagnostic for Bayesian hierarchical models
- Frequentist variability of Bayesian posteriors with the infinitesimal jackknife
 - Use sensitivity of mean parameters to speed up bootstrapping MCMC
- Flow-based Bayesian local robustness: “Otto grad”
 - Use easy-to-compute approximate flows to speed up bootstrapping MCMC

- Peter Bühlmann. Invariance, Causality and Robustness: 2018 Neyman Lecture. 35(3):404–426. ISSN 0883-4237. URL <https://www.jstor.org/stable/26997912>.
- Rafael Cabral, David Bolin, and Håvard Rue. Robustness, model checking, and hierarchical models. 87(3):632–652. ISSN 1369-7412, 1467-9868. doi: 10.1093/jrsssb/qkae107. URL <https://academic.oup.com/jrsssb/article/87/3/632/7914952>.
- R. Giordano and T. Broderick. The Bayesian infinitesimal jackknife for variance. *arXiv preprint arXiv:2305.06466*, 2026.
- R. Giordano, T. Broderick, and M. Jordan. Covariances, robustness and variational Bayes. *The Journal of Machine Learning Research*, 19(1):1981–2029, 2018.
- Sven Kuh, Lauren Kennedy, Qixuan Chen, and Andrew Gelman. Using leave-one-out cross validation (LOO) in a multilevel regression and poststratification (MRP) workflow: A cautionary tale. 43(5):953–982. ISSN 1097-0258. doi: 10.1002/sim.9964. URL <https://onlinelibrary.wiley.com/doi/abs/10.1002/sim.9964>.
- Krikamol Muandet, Wittawat Jitkrittum, and Jonas Kübler. Kernel Conditional Moment Test via Maximum Moment Restriction. In *Proceedings of the 36th Conference on Uncertainty in Artificial Intelligence (UAI)*, pages 41–50. PMLR. URL <https://proceedings.mlr.press/v124/muandet20a.html>.
- Mateo Rojas-Carulla, Bernhard Scholkopf, Richard Turner, and Jonas Peters. Invariant Models for Causal Transfer Learning.
- Aki Vehtari, Andrew Gelman, and Jonah Gabry. Practical Bayesian model evaluation using leave-one-out cross-validation and WAIC. 27(5):1413–1432. ISSN 1573-1375. doi: 10.1007/s11222-016-9696-4. URL <https://doi.org/10.1007/s11222-016-9696-4>.