

Next-generation Bayesian local robustness

Ryan Giordano (rgiordano@berkeley.edu, UC Berkeley)

Berkeley Neyman Seminar


September 2026

The basic problem


We have a survey population, for whom we observe:

- Covariates x (e.g. race, gender, zip code, age, education level)
- Responses y (e.g. A binary response to “do you support candidate Z”)

We want the average response in a target population, in which we observe only covariates.



Observe (x_i, y_i) for $i = 1, \dots, N_S$
(This is the “Data”, $\mathcal{D}$)



Observe x_j for $j = 1, \dots, N_T$

The basic problem

We have a survey population, for whom we observe:


- Covariates x (e.g. race, gender, zip code, age, education level)
- Responses y (e.g. A binary response to “do you support candidate Z”)

We want the average response in a target population, in which we observe only covariates.

The problem is that it may be that $\mathcal{P}_S(x) \neq \mathcal{P}_T(x)$. (Assume $\mathcal{P}(y|x)$ the same.)



Observe (x_i, y_i) for $i = 1, \dots, N_S$
(This is the “Data”, $\mathcal{D}$)



Observe x_j for $j = 1, \dots, N_T$

The basic problem

We have a survey population, for whom we observe:


- Covariates x (e.g. race, gender, zip code, age, education level)
- Responses y (e.g. A binary response to “do you support candidate Z”)

We want the average response in a target population, in which we observe only covariates.

The problem is that it may be that $\mathcal{P}_S(x) \neq \mathcal{P}_T(x)$. (Assume $\mathcal{P}(y|x)$ the same.)



Observe (x_i, y_i) for $i = 1, \dots, N_S$
(This is the “Data”, $\mathcal{D}$)



Observe x_j for $j = 1, \dots, N_T$

This talk will study “Multilevel regression and poststratification” estimators [Gel97]:

- Set $\mathbb{E}[y|x, \theta] = m(\theta^\top x)$ and prior $\mathcal{P}(\theta)$ (e.g. Hierarchical logistic regression)
- Estimate the posterior $\mathcal{P}(\theta|\mathcal{D})$ with Markov Chain Monte Carlo (MCMC)
- Take $\hat{\mu} = \frac{1}{N_T} \sum_{j=1}^{N_T} \hat{y}_j$ where $\hat{y}_j = \mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})}[y|x_j]$.

The basic problem

We have a survey population, for whom we observe:


- Covariates x (e.g. race, gender, zip code, age, education level)
- Responses y (e.g. A binary response to “do you support candidate Z”)

We want the average response in a target population, in which we observe only covariates.

The problem is that it may be that $\mathcal{P}_S(x) \neq \mathcal{P}_T(x)$. (Assume $\mathcal{P}(y|x)$ the same.)



Observe (x_i, y_i) for $i = 1, \dots, N_S$
(This is the “Data”, $\mathcal{D}$)



Observe x_j for $j = 1, \dots, N_T$

This talk will study “Multilevel regression and poststratification” estimators [Gel97]:

- Set $\mathbb{E}[y|x, \theta] = m(\theta^\top x)$ and prior $\mathcal{P}(\theta)$ (e.g. Hierarchical logistic regression)
- Estimate the posterior $\mathcal{P}(\theta|\mathcal{D})$ with Markov Chain Monte Carlo (MCMC)
- Take $\hat{\mu} = \frac{1}{N_T} \sum_{j=1}^{N_T} \hat{y}_j$ where $\hat{y}_j = \mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})}[y|x_j]$.

Problem: MCMC is time consuming and the regression model is surely misspecified.

Goal: Meaningful regression specification checks with frequentist uncertainty.

Result preview




Figure 1: A covariate shift diagnostic

Here is our check for two models on the same simulated data: one correctly specified and one misspecified. The misspecified model is more extreme.

Problem: What is “large enough” to be real evidence of misspecification?

Possible answer: Frequentist confidence intervals via the parametric bootstrap.

Result preview




Figure 1: A covariate shift diagnostic

The frequentist interval based on the parametric bootstrap rejects the misspecified model but not the correctly specified model.

Problem: Expensive: each bootstrap sample required a new MCMC run.

Possible answer: Infinitesimal jackknife confidence intervals? [[GB26](#)]

Result preview




Figure 1: A covariate shift diagnostic

IJ confidence intervals are fast and often work well for posterior means.
But they are not accurate for this more complex diagnostic (high Monte Carlo error).

Result preview




Figure 1: A covariate shift diagnostic

IJ confidence intervals are fast and often work well for posterior means.
But they are not accurate for this more complex diagnostic (high Monte Carlo error).

Answer: Local flow sensitivity measures produce good approximations to the bootstrap at negligible computational expense.

- Local checks of specification of Bayesian regression models
 - (With Alice Cima, Lucas Schwengber, Erin Hartman, Jared Miller, Avi Feller)
- Frequentist variability of Bayesian posteriors with the infinitesimal jackknife
 - (With Tamara Broderick)
- Bayesian robustness with local flows
 - (With Lucas Schwengber)



Local checks of specification of Bayesian regression models

$$\hat{\mu} = \frac{1}{N_T} \sum_{j=1}^{N_T} \hat{y}_j = \mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})} \underbrace{\left[\overbrace{\frac{1}{N_T} \sum_{j=1}^{N_T} m(\theta^\top x_j)}^{\text{Just a function } \phi(\theta)} \right]}_{\text{Want to check specification without re-running MCMC}}.$$

Some existing local approaches:

- Use out-of-sample checking [VGG17; Gel+20; Kuh+24]
 - The task is not to predict on a new IID survey datapoint, but to predict in a new population
- Specify an alternative model, maybe test with Bayes factor [DRM25; CBR25]
 - Bayes factors can depend strongly on priors (Lindley's paradox)
 - Articulating a family of models is inconvenient
 - Doesn't take the target into account

Local specification checks

$$\hat{\mu} = \frac{1}{N_T} \sum_{j=1}^{N_T} \hat{y}_j = \mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})} \underbrace{\left[\overbrace{\frac{1}{N_T} \sum_{j=1}^{N_T} m(\theta^\top x_j)}^{\text{Just a function } \phi(\theta)} \right]}_{\substack{\text{Want to check specification} \\ \text{without re-running MCMC}}}.$$

Some existing local approaches:

- Use out-of-sample checking [VGG17; Gel+20; Kuh+24]
 - The task is not to predict on a new IID survey datapoint, but to predict in a new population
- Specify an alternative model, maybe test with Bayes factor [DRM25; CBR25]
 - Bayes factors can depend strongly on priors (Lindley's paradox)
 - Articulating a family of models is inconvenient
 - Doesn't take the target into account

The model is definitely misspecified, we want to check whether it's *importantly* misspecified.

We will check the “stability” component in the PCS framework [YB24].

Local specification checks

$$\hat{\mu} = \frac{1}{N_T} \sum_{j=1}^{N_T} \hat{y}_j = \underbrace{\mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})} \left[\frac{1}{N_T} \sum_{j=1}^{N_T} m(\theta^\top x_j) \right]}_{\text{Want to check specification without re-running MCMC}}.$$

Why do we care about correct specification? We want to use $\mathcal{P}_S(x)$ to learn about $\mathcal{P}_T(x)$.

Correct specification $\Rightarrow$ Regression estimates are invariant to $\mathcal{P}_S(x)$

Local specification checks

$$\hat{\mu} = \frac{1}{N_T} \sum_{j=1}^{N_T} \hat{y}_j = \underbrace{\mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})} \left[\frac{1}{N_T} \sum_{j=1}^{N_T} m(\theta^\top x_j) \right]}_{\substack{\text{Want to check specification} \\ \text{without re-running MCMC}}}.$$

Why do we care about correct specification? We want to use $\mathcal{P}_S(x)$ to learn about $\mathcal{P}_T(x)$.

Correct specification $\Rightarrow$ Regression estimates are invariant to $\mathcal{P}_S(x)$
This is what we actually want — let's check this directly.

Standard idea in causal, kernel learning, transfer learning [Shi00; Böh20; Roj+; MJK20].

$$\hat{\mu} = \frac{1}{N_T} \sum_{j=1}^{N_T} \hat{y}_j = \underbrace{\mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})} \left[\frac{1}{N_T} \sum_{j=1}^{N_T} m(\theta^\top x_j) \right]}_{\substack{\text{Want to check specification} \\ \text{without re-running MCMC}}}.$$

Why do we care about correct specification? We want to use $\mathcal{P}_S(x)$ to learn about $\mathcal{P}_T(x)$.

Correct specification $\Rightarrow$ Regression estimates are invariant to $\mathcal{P}_S(x)$
This is what we actually want — let's check this directly.

Standard idea in causal, kernel learning, transfer learning [Shi00; Böh20; Roj+; MJK20].

We provide a local Bayesian diagnostic version of this idea:

1. Let $\mathcal{P}_{\mathcal{D}}(x, y)$ denote a generic joint density of x, y .
2. Write our estimand as a functional $\hat{\mu}(\mathcal{P}_{\mathcal{D}})$
3. Define an additive covariate density perturbation $w(x)$
4. Define $\mathcal{P}_{\mathcal{D}}^\epsilon(x, y) = (\mathcal{P}_{\mathcal{D}}(x) + \epsilon w(x)) \mathcal{P}_{\mathcal{D}}(y|x)$
5. Compute the derivative $\Delta(\mathcal{D}) = \partial \hat{\mu}(\mathcal{P}_{\mathcal{D}}^\epsilon) / \partial \epsilon$ and check whether ≈ 0

Simple visual example




Figure 2: A covariate shift diagnostic

Consider logistic regression with two regressors, x_1 and x_2 .

Simple visual example




Figure 2: A covariate shift diagnostic

Consider logistic regression with two regressors, x_1 and x_2 .

Taking $w(x) = \mathbb{I}(x_1 > 1)$, we tend to upweight more positive examples.

Simple visual example




Figure 2: A covariate shift diagnostic

Consider logistic regression with two regressors, x_1 and x_2 .

Taking $w(x) = \mathbb{I}(x_1 > 1)$, we tend to upweight more positive examples.

With enough data, a regression that omits x_1 will change its coefficient of x_2 under re-weighting, whereas a correctly specified regression will not need to.

Bayesian sensitivity analysis

Let the regression log likelihood be $\ell(y|x, \theta)$, and fix $w(x)$. Then define

$$\mathcal{P}(\theta|\mathcal{D}) \propto \exp\left(\sum_{i=1}^{N_S} \ell(y_i|x_i, \theta)\right) \pi(\theta) \text{ and } \mathcal{P}^\epsilon(\theta) \propto \mathcal{P}(\theta|\mathcal{D}) \exp\left(\sum_{i=1}^{N_S} \epsilon w(x_i) \ell(y_i|x_i, \theta)\right)$$

Bayesian sensitivity analysis

Let the regression log likelihood be $\ell(y|x, \theta)$, and fix $w(x)$. Then define

$$\mathcal{P}(\theta|\mathcal{D}) \propto \exp\left(\sum_{i=1}^{N_S} \ell(y_i|x_i, \theta)\right) \pi(\theta) \text{ and } \mathcal{P}^\epsilon(\theta) \propto \mathcal{P}(\theta|\mathcal{D}) \exp\left(\sum_{i=1}^{N_S} \epsilon w(x_i) \ell(y_i|x_i, \theta)\right)$$

Theorem: [DF86; Gus96; [GBJ18](#)]:

Consider a family of distributions $\mathcal{P}^\epsilon(\theta) \propto \mathcal{P}(\theta) \exp(\epsilon h(\theta))$ and a function $\phi(\theta)$.

By interchange of integration and differentiation,

$$\left. \frac{\partial \mathbb{E}_{\mathcal{P}^\epsilon(\theta)} [\phi(\theta)]}{\partial \epsilon} \right|_{\epsilon=0} = \text{Cov}_{\mathcal{P}(\theta)} (\phi(\theta), h(\theta)) .$$

Bayesian sensitivity analysis

Let the regression log likelihood be $\ell(y|x, \theta)$, and fix $w(x)$. Then define

$$\mathcal{P}(\theta|\mathcal{D}) \propto \exp\left(\sum_{i=1}^{N_S} \ell(y_i|x_i, \theta)\right) \pi(\theta) \text{ and } \mathcal{P}^\epsilon(\theta) \propto \mathcal{P}(\theta|\mathcal{D}) \exp\left(\sum_{i=1}^{N_S} \epsilon w(x_i) \ell(y_i|x_i, \theta)\right)$$

Theorem: [DF86; Gus96; [GBJ18](#)]:

Consider a family of distributions $\mathcal{P}^\epsilon(\theta) \propto \mathcal{P}(\theta) \exp(\epsilon h(\theta))$ and a function $\phi(\theta)$.

By interchange of integration and differentiation,

$$\left. \frac{\partial \mathbb{E}_{\mathcal{P}^\epsilon(\theta)} [\phi(\theta)]}{\partial \epsilon} \right|_{\epsilon=0} = \text{Cov}_{\mathcal{P}(\theta)} (\phi(\theta), h(\theta)) .$$

Our diagnostic is then given by

$$\Delta(\mathcal{D}) := \left. \frac{\partial \hat{\mu}}{\partial \epsilon} \right|_{\epsilon=0} = \text{Cov}_{\mathcal{P}(\theta|\mathcal{D})} \left(\underbrace{\frac{1}{N_T} \sum_{j=1}^{N_T} m(\theta^\top x_j)}_{\text{Estimand}}, \underbrace{\sum_{i=1}^{N_S} w(x_i) \ell(y_i|x_i, \theta)}_{\text{Log model perturbation}} \right)$$

We expect $\Delta(\mathcal{D})$ to be *close* to zero when the model is correctly specified.

Bootstrapping Bayesian sensitivity analysis

We expect $\Delta(\mathcal{D})$ to be *close* to zero when the model is correctly specified.

But on actual data, $\Delta(\mathcal{D})$ will not be precisely zero, even under correct specification.

Bootstrapping Bayesian sensitivity analysis

We expect $\Delta(\mathcal{D})$ to be *close* to zero when the model is correctly specified.

But on actual data, $\Delta(\mathcal{D})$ will not be precisely zero, even under correct specification.




Figure 3: A covariate shift diagnostic

Simulation data: Logistic regression with five active regressors, each with covariate shift.

Correctly specified model: $y \sim \text{Logistic}(1 + X1 + X2 + X3 + X4 + X5)$

Misspecified model: $y \sim \text{Logistic}(1 + X3 + X4 + X5)$

Diagnostic uses $w(x) = \mathbb{I}(X2 > 0.878)$

Bootstrapping Bayesian sensitivity analysis

We expect $\Delta(\mathcal{D})$ to be *close* to zero when the model is correctly specified.

But on actual data, $\Delta(\mathcal{D})$ will not be precisely zero, even under correct specification.

To assess “closeness to zero,” compare with $\Delta(\mathcal{D}^*)$, where $\mathcal{D}^*$ is a bootstrap sample.

- Compute $\Delta(\mathcal{D})$ on the original dataset
- For $b = 1, \dots, 100$:
 1. Sample $\mathcal{D}_b^* := (x_1^*, y_1^*), \dots (x_{N_S}^*, y_{N_S}^*)$ (parametric or nonparametric bootstrap)
 2. Run MCMC to estimate $\mathcal{P}(\theta|\mathcal{D}_b^*)$ (**time consuming!**)
 3. Use the draws to estimate $\Delta(\mathcal{D}_b^*)$
- Compare $\Delta(\mathcal{D})$ with distribution of $\Delta(\mathcal{D}_b^*)$.

Bootstrapping Bayesian sensitivity analysis

We expect $\Delta(\mathcal{D})$ to be *close* to zero when the model is correctly specified.

But on actual data, $\Delta(\mathcal{D})$ will not be precisely zero, even under correct specification.




Figure 3: A covariate shift diagnostic

Simulation data: Logistic regression with five active regressors, each with covariate shift.

Correctly specified model: $y \sim \text{Logistic}(1 + X1 + X2 + X3 + X4 + X5)$

Misspecified model: $y \sim \text{Logistic}(1 + X3 + X4 + X5)$

Diagnostic uses $w(x) = \mathbb{I}(X2 > 0.878)$

Frequentist variance with the infinitesimal
jackknife

Why frequentist variance?

For the remainder of the talk, focus on *fast approximations* to the *frequentist* sampling distribution under $x_i, y_i \stackrel{iid}{\sim} \mathcal{P}_S(x, y)$ of Bayesian posterior quantities like our diagnostic.

We will form *local approximations* computable only from *the original posterior*.

Why frequentist variance?

For the remainder of the talk, focus on *fast approximations* to the *frequentist* sampling distribution under $x_i, y_i \stackrel{iid}{\sim} \mathcal{P}_S(x, y)$ of Bayesian posterior quantities like our diagnostic.

We will form *local approximations* computable only from *the original posterior*.

No need to commit to our covariate shift diagnostic to care!

- Frequentist variability of *any* model checking diagnostic is interesting
- Frequentist variability meaningful under misspecification [HM23; WDH26; LJ26; JLR25]

Data re-weighting (again)

Let the regression log likelihood be $\ell(y|x, \theta)$, and fix $w(x)$. Then define

$$\mathcal{P}(\theta|\mathcal{D}) \propto \exp\left(\sum_{i=1}^{N_S} \ell(y_i|x_i, \theta)\right) \pi(\theta) \text{ and } \mathcal{P}^\epsilon(\theta) \propto \mathcal{P}(\theta|\mathcal{D}) \exp\left(\sum_{i=1}^{N_S} \epsilon w(x_i) \ell(y_i|x_i, \theta)\right)$$

Data re-weighting (again)

Let the regression log likelihood be $\ell(y|x, \theta)$, and fix $w(x)$. Then define

$$\mathcal{P}(\theta|\mathcal{D}) \propto \exp\left(\sum_{i=1}^{N_S} \ell(y_i|x_i, \theta)\right) \pi(\theta) \text{ and } \mathcal{P}^\epsilon(\theta) \propto \mathcal{P}(\theta|\mathcal{D}) \exp\left(\sum_{i=1}^{N_S} \epsilon w(x_i) \ell(y_i|x_i, \theta)\right)$$

If we let $w(x_i)$ denote a *bootstrap weight*, then $\mathcal{P}^\epsilon$ at $\epsilon = 1$ is the bootstrapped posterior!

Data re-weighting (again)

Let the regression log likelihood be $\ell(y|x, \theta)$, and fix $w(x)$. Then define

$$\mathcal{P}(\theta|\mathcal{D}) \propto \exp \left(\sum_{i=1}^{N_S} \ell(y_i|x_i, \theta) \right) \pi(\theta) \text{ and } \mathcal{P}^\epsilon(\theta) \propto \mathcal{P}(\theta|\mathcal{D}) \exp \left(\sum_{i=1}^{N_S} \epsilon w(x_i) \ell(y_i|x_i, \theta) \right)$$

If we let $w(x_i)$ denote a *bootstrap weight*, then $\mathcal{P}^\epsilon$ at $\epsilon = 1$ is the bootstrapped posterior!

We can form a series approximation:

$$\mathbb{E}_{\mathcal{P}(\theta|\mathcal{D}^*)} [\phi(\theta)] \approx \mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})} [\phi(\theta)] + \left. \frac{\partial \mathbb{E}_{\mathcal{P}^\epsilon(\theta)} [\phi(\theta)]}{\partial \epsilon} \right|_{\epsilon=0} (1 - 0)$$

Data re-weighting (again)

Let the regression log likelihood be $\ell(y|x, \theta)$, and fix $w(x)$. Then define

$$\mathcal{P}(\theta|\mathcal{D}) \propto \exp \left(\sum_{i=1}^{N_S} \ell(y_i|x_i, \theta) \right) \pi(\theta) \text{ and } \mathcal{P}^\epsilon(\theta) \propto \mathcal{P}(\theta|\mathcal{D}) \exp \left(\sum_{i=1}^{N_S} \epsilon w(x_i) \ell(y_i|x_i, \theta) \right)$$

If we let $w(x_i)$ denote a *bootstrap weight*, then $\mathcal{P}^\epsilon$ at $\epsilon = 1$ is the bootstrapped posterior!

We can form a series approximation:

$$\begin{aligned} \mathbb{E}_{\mathcal{P}(\theta|\mathcal{D}^*)} [\phi(\theta)] &\approx \mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})} [\phi(\theta)] + \frac{\partial \mathbb{E}_{\mathcal{P}^\epsilon(\theta)} [\phi(\theta)]}{\partial \epsilon} \bigg|_{\epsilon=0} (1 - 0) \\ &= \mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})} [\phi(\theta)] + \underbrace{\sum_{i=1}^{N_S} \frac{\partial \mathbb{E}_{\mathcal{P}^\epsilon(\theta)} [\phi(\theta)]}{\partial w_i} \bigg|_{w_i=0}}_{\text{"Empirical influence function"} \frac{\psi_i}{N_S}} w_i \end{aligned}$$

Data re-weighting (again)

Let the regression log likelihood be $\ell(y|x, \theta)$, and fix $w(x)$. Then define

$$\mathcal{P}(\theta|\mathcal{D}) \propto \exp \left(\sum_{i=1}^{N_S} \ell(y_i|x_i, \theta) \right) \pi(\theta) \text{ and } \mathcal{P}^\epsilon(\theta) \propto \mathcal{P}(\theta|\mathcal{D}) \exp \left(\sum_{i=1}^{N_S} \epsilon w(x_i) \ell(y_i|x_i, \theta) \right)$$

If we let $w(x_i)$ denote a *bootstrap weight*, then $\mathcal{P}^\epsilon$ at $\epsilon = 1$ is the bootstrapped posterior!

We can form a series approximation:

$$\begin{aligned} \mathbb{E}_{\mathcal{P}(\theta|\mathcal{D}^*)} [\phi(\theta)] &\approx \mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})} [\phi(\theta)] + \frac{\partial \mathbb{E}_{\mathcal{P}^\epsilon(\theta)} [\phi(\theta)]}{\partial \epsilon} \bigg|_{\epsilon=0} (1 - 0) \\ &= \mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})} [\phi(\theta)] + \underbrace{\sum_{i=1}^{N_S} \frac{\partial \mathbb{E}_{\mathcal{P}^\epsilon(\theta)} [\phi(\theta)]}{\partial w_i} \bigg|_{w_i=0}}_{\text{"Empirical influence function"} \frac{\psi_i}{N_S}} w_i \\ &= \mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})} [\phi(\theta)] + \frac{1}{N_S} \sum_{i=1}^{N_S} \psi_i w_i \end{aligned}$$

Data re-weighting (again)

Let the regression log likelihood be $\ell(y|x, \theta)$, and fix $w(x)$. Then define

$$\mathcal{P}(\theta|\mathcal{D}) \propto \exp\left(\sum_{i=1}^{N_S} \ell(y_i|x_i, \theta)\right) \pi(\theta) \text{ and } \mathcal{P}^\epsilon(\theta) \propto \mathcal{P}(\theta|\mathcal{D}) \exp\left(\sum_{i=1}^{N_S} \epsilon w(x_i) \ell(y_i|x_i, \theta)\right)$$

If we let $w(x_i)$ denote a *bootstrap weight*, then $\mathcal{P}^\epsilon$ at $\epsilon = 1$ is the bootstrapped posterior!

We can form a series approximation:

$$\begin{aligned} \mathbb{E}_{\mathcal{P}(\theta|\mathcal{D}^*)} [\phi(\theta)] &\approx \mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})} [\phi(\theta)] + \left. \frac{\partial \mathbb{E}_{\mathcal{P}^\epsilon(\theta)} [\phi(\theta)]}{\partial \epsilon} \right|_{\epsilon=0} (1 - 0) \\ &= \mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})} [\phi(\theta)] + \underbrace{\sum_{i=1}^{N_S} \left. \frac{\partial \mathbb{E}_{\mathcal{P}^\epsilon(\theta)} [\phi(\theta)]}{\partial w_i} \right|_{w_i=0}}_{\text{"Empirical influence function"} \frac{\psi_i}{N_S}} w_i \\ &= \mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})} [\phi(\theta)] + \frac{1}{N_S} \sum_{i=1}^{N_S} \psi_i w_i \end{aligned}$$

If this series is accurate, then we can compute ψ_i using the original posterior and quickly approximate bootstrap draws.

Task: Prove that the series is asymptotically accurate enough for hypothesis testing.

Bayesian von-Mises Expansion

Define the “generalized posterior” functional [GB26]:

$$T(\mathcal{P}, N) := \frac{\int \phi(\theta) \exp \left(N \int \ell(y|x, \theta) \mathcal{P}(x, y) dx dy \right) \pi(\theta) d\theta}{\int \exp \left(N \int \ell(y|x, \theta) \mathcal{P}(x, y) dx dy \right) \pi(\theta) d\theta}.$$

Note that $\mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})} [\phi(\theta)] = T(\hat{\mathcal{P}}_S, N_S)$

Bayesian von-Mises Expansion

Define the “generalized posterior” functional [GB26]:

$$T(\mathcal{P}, N) := \frac{\int \phi(\theta) \exp \left(N \int \ell(y|x, \theta) \mathcal{P}(x, y) dx dy \right) \pi(\theta) d\theta}{\int \exp \left(N \int \ell(y|x, \theta) \mathcal{P}(x, y) dx dy \right) \pi(\theta) d\theta}.$$

Note that $\mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})} [\phi(\theta)] = T(\hat{\mathcal{P}}_S, N_S)$

To derive frequentist variance estimates, we study the *von Mises expansion* [von47].

Let $\tilde{\mathcal{P}}^t = t\hat{\mathcal{P}}_S + (1-t)\mathcal{P}_S$. Then form the series expansion in t :

$$\begin{aligned} \sqrt{N_S} \left(\mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})} [\phi(\theta)] - T(\mathcal{P}_S, N_S) \right) &= \sqrt{N_S} \left(T(\hat{\mathcal{P}}_S, N_S) - T(\mathcal{P}_S, N_S) \right) \\ &= \sqrt{N_S} \left. \frac{\partial T(\tilde{\mathcal{P}}^t, N_S)}{\partial t} \right|_{t=0} (1-0) + \mathcal{E}(\tilde{t}) \quad (\tilde{t} \in [0, 1]) \\ &= \sqrt{N_S} \underbrace{\sum_{n=1}^N (\psi_i - \bar{\psi})}_{\text{Infinitesimal jackknife estimator}} + o_p(1) + \mathcal{E}(\tilde{t}). \end{aligned}$$

Bayesian von-Mises Expansion Results

$$\sqrt{N_S} \left(\mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})} [\phi(\theta)] - T(\mathcal{P}_S, N_S) \right) = \sqrt{N_S} \sum_{n=1}^N (\psi_i - \bar{\psi}) + o_p(1) + \mathcal{E}(\tilde{t}).$$

Theorem: [GB26] (sketch):

Under Bernstein-von Mises-like conditions, $\sup_{\tilde{t} \in [0,1]} |\mathcal{E}(\tilde{t})| \xrightarrow{P} 0$.

Corollary: [von47] Let $\frac{1}{N_S} \sum_{i=1}^{N_S} (\psi_i - \bar{\psi})^2 \xrightarrow{P} V$. By the CLT applied to ψ_i ,

$$\sqrt{N_S} \left(\mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})} [\phi(\theta)] - T(\mathcal{P}_S, N_S) \right) \xrightarrow{d} \mathcal{N}(0, V).$$

Bayesian von-Mises Expansion Results

$$\sqrt{N_S} \left(\mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})} [\phi(\theta)] - T(\mathcal{P}_S, N_S) \right) = \sqrt{N_S} \sum_{n=1}^N (\psi_i - \bar{\psi}) + o_p(1) + \mathcal{E}(\tilde{t}).$$

Theorem: [GB26] (sketch):

Under Bernstein-von Mises-like conditions, $\sup_{\tilde{t} \in [0,1]} |\mathcal{E}(\tilde{t})| \xrightarrow{P} 0$.

Corollary: [von47] Let $\frac{1}{N_S} \sum_{i=1}^{N_S} (\psi_i - \bar{\psi})^2 \xrightarrow{P} V$. By the CLT applied to ψ_i ,

$$\sqrt{N_S} \left(\mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})} [\phi(\theta)] - T(\mathcal{P}_S, N_S) \right) \xrightarrow{d} \mathcal{N}(0, V).$$

To compute and study the parametric bootstrap instead is a simple change:

$$\mathcal{P}^\epsilon(\theta) \propto \mathcal{P}(\theta|\mathcal{D}) \exp \left(\sum_{i=1}^{N_S} \epsilon (\ell(y_i^*|x_i, \theta) - \ell(y_i|x_i, \theta)) \right).$$

(See, e.g., [Giordano et al. \[Gio+26\]](#) Theorem 4.1).

Bayesian von-Mises Expansion Results

$$\sqrt{N_S} \left(\mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})} [\phi(\theta)] - T(\mathcal{P}_S, N_S) \right) = \sqrt{N_S} \sum_{n=1}^N (\psi_i - \bar{\psi}) + o_p(1) + \mathcal{E}(\tilde{t}).$$

Theorem: [GB26] (sketch):

Under Bernstein-von Mises-like conditions, $\sup_{\tilde{t} \in [0,1]} |\mathcal{E}(\tilde{t})| \xrightarrow{P} 0$.

Corollary: [von47] Let $\frac{1}{N_S} \sum_{i=1}^{N_S} (\psi_i - \bar{\psi})^2 \xrightarrow{P} V$. By the CLT applied to ψ_i ,

$$\sqrt{N_S} \left(\mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})} [\phi(\theta)] - T(\mathcal{P}_S, N_S) \right) \xrightarrow{d} \mathcal{N}(0, V).$$

To compute and study the parametric bootstrap instead is a simple change:

$$\mathcal{P}^\epsilon(\theta) \propto \mathcal{P}(\theta|\mathcal{D}) \exp \left(\sum_{i=1}^{N_S} \epsilon (\ell(y_i^*|x_i, \theta) - \ell(y_i|x_i, \theta)) \right).$$

(See, e.g., [Giordano et al. \[Gio+26\]](#) Theorem 4.1).

Note: The von-Mises expansion is stronger than necessary for consistency [GB26]. But it can be used to prove *lower bounds* on $\mathcal{E}(\tilde{t})$ and so *inconsistency* when no BVM theorem holds.




Figure 4: MrP Posterior Mean

The IJ does great for the MrP point estimate.

See also Luo and Ji [LJ26] and Ji, Lee, and Rabe-Hesketh [JLR25].




Figure 4: MrP Posterior Variance

The IJ does pretty well for the MrP posterior variance.




Figure 4: Covariate Shift Diagnostic

Need something better than the IJ for noisier / more complex posterior quantities!

Bayesian robustness with local flows

Why flows for sensitivity analysis?

Current Bayesian local robustness is built around linearizing posterior expectations:

$$\mathbb{E}_{\mathcal{P}^\epsilon(\theta)} [\phi(\theta)] \approx \mathbb{E}_{\mathcal{P}(\theta)} [\phi(\theta)] + \left. \frac{\partial \mathbb{E}_{\mathcal{P}^\epsilon(\theta)} [\phi(\theta)]}{\partial \epsilon} \right|_{\epsilon=0} (\epsilon - 0),$$

where the derivative is estimated with MCMC [Gus96; [GBJ18](#); CBR25; DRM25; [Gio+26](#)].

Why flows for sensitivity analysis?

Current Bayesian local robustness is built around linearizing posterior expectations:

$$\mathbb{E}_{\mathcal{P}^{\epsilon}(\theta)} [\phi(\theta)] \approx \mathbb{E}_{\mathcal{P}(\theta)} [\phi(\theta)] + \left. \frac{\partial \mathbb{E}_{\mathcal{P}^{\epsilon}(\theta)} [\phi(\theta)]}{\partial \epsilon} \right|_{\epsilon=0} (\epsilon - 0),$$

where the derivative is estimated with MCMC [Gus96; [GBJ18](#); CBR25; DRM25; [Gio+26](#)].

This has some shortcomings:

- Potential MCMC variance, not obvious how to smooth
- Linear extrapolation can lead to impossible values (e.g. negative variances)
- Complex quantities like posterior quantiles are hard to linearize
- Bayesians like to work with draws, not analytic expectations [HGV25]

Why flows for sensitivity analysis?

Current Bayesian local robustness is built around linearizing posterior expectations:

$$\mathbb{E}_{\mathcal{P}^\epsilon(\theta)} [\phi(\theta)] \approx \mathbb{E}_{\mathcal{P}(\theta)} [\phi(\theta)] + \left. \frac{\partial \mathbb{E}_{\mathcal{P}^\epsilon(\theta)} [\phi(\theta)]}{\partial \epsilon} \right|_{\epsilon=0} (\epsilon - 0),$$

where the derivative is estimated with MCMC [Gus96; GBJ18; CBR25; DRM25; Gio+26].

This has some shortcomings:

- Potential MCMC variance, not obvious how to smooth
- Linear extrapolation can lead to impossible values (e.g. negative variances)
- Complex quantities like posterior quantiles are hard to linearize
- Bayesians like to work with draws, not analytic expectations [HGV25]

Idea: Instead of linearizing expectations, learn a flow from $\mathcal{P}$ to $\mathcal{P}^\epsilon$.

There's a large literature on learning flows between distant distributions. [Pap+21]

Our problem is easier than much of this literature:

- Only need “local” flows
- Our perturbations have a lot of structure (e.g. bootstraps)

Consider $\mathcal{P}^\epsilon(\theta) \propto \mathcal{P}(\theta) \exp(\epsilon h(\theta))$, with draws $\theta \sim \mathcal{P}(\theta)$.




Figure 5: Draws from a posterior

Consider $\mathcal{P}^\epsilon(\theta) \propto \mathcal{P}(\theta) \exp(\epsilon h(\theta))$, with draws $\theta \sim \mathcal{P}(\theta)$.




Figure 5: A posterior perturbation

Imagine the density $\mathcal{P}(\theta)$ coming out of the page. The perturbation $h(\theta)$ is defined “vertically.”

Consider $\mathcal{P}^\epsilon(\theta) \propto \mathcal{P}(\theta) \exp(\epsilon h(\theta))$, with draws $\theta \sim \mathcal{P}(\theta)$.




Figure 5: A posterior perturbation with weights

Plug-in estimates for the covariance are equivalent to producing the perturbation by weighting [\[GBJ18, Appendix B\]](#)

Consider $\mathcal{P}^\epsilon(\theta) \propto \mathcal{P}(\theta) \exp(\epsilon h(\theta))$, with draws $\theta \sim \mathcal{P}(\theta)$.




Figure 5: A posterior perturbation with a flow

We will try, instead, try to effect the perturbation with a flow $\theta \mapsto G^\epsilon(\theta)$

Consider $\mathcal{P}^\epsilon(\theta) \propto \mathcal{P}(\theta) \exp(\epsilon h(\theta))$.

We want a flow $\theta \mapsto G^\epsilon(\theta)$ such that, for small ϵ , $G^\epsilon(\theta) \sim \mathcal{P}^\epsilon(\theta)$ and $G^0(\theta) = \theta \sim \mathcal{P}(\theta)$.

Consider $\mathcal{P}^\epsilon(\theta) \propto \mathcal{P}(\theta) \exp(\epsilon h(\theta))$.

We want a flow $\theta \mapsto G^\epsilon(\theta)$ such that, for small ϵ , $G^\epsilon(\theta) \sim \mathcal{P}^\epsilon(\theta)$ and $G^0(\theta) = \theta \sim \mathcal{P}(\theta)$.

Suppose there is a scalar-valued potential function $u(\theta)$ such that

$$G^\epsilon(\theta) = \theta + \epsilon \nabla u(\theta) + O(\epsilon^2) \quad \text{as } \epsilon \rightarrow 0. \text{ Then:}$$

Consider $\mathcal{P}^\epsilon(\theta) \propto \mathcal{P}(\theta) \exp(\epsilon h(\theta))$.

We want a flow $\theta \mapsto G^\epsilon(\theta)$ such that, for small ϵ , $G^\epsilon(\theta) \sim \mathcal{P}^\epsilon(\theta)$ and $G^0(\theta) = \theta \sim \mathcal{P}(\theta)$.

Suppose there is a scalar-valued potential function $u(\theta)$ such that

$$G^\epsilon(\theta) = \theta + \epsilon \nabla u(\theta) + O(\epsilon^2) \quad \text{as } \epsilon \rightarrow 0. \text{ Then:}$$

$$\mathbb{E}_{\mathcal{P}^\epsilon(\theta)} [\phi(\theta)] = \mathbb{E}_{\mathcal{P}(\theta)} [\phi(G^\epsilon(\theta))] \Rightarrow \quad \text{(good flow)}$$

Consider $\mathcal{P}^\epsilon(\theta) \propto \mathcal{P}(\theta) \exp(\epsilon h(\theta))$.

We want a flow $\theta \mapsto G^\epsilon(\theta)$ such that, for small ϵ , $G^\epsilon(\theta) \sim \mathcal{P}^\epsilon(\theta)$ and $G^0(\theta) = \theta \sim \mathcal{P}(\theta)$.

Suppose there is a scalar-valued potential function $u(\theta)$ such that

$$G^\epsilon(\theta) = \theta + \epsilon \nabla u(\theta) + O(\epsilon^2) \quad \text{as } \epsilon \rightarrow 0. \text{ Then:}$$

$$\mathbb{E}_{\mathcal{P}^\epsilon(\theta)} [\phi(\theta)] = \mathbb{E}_{\mathcal{P}(\theta)} [\phi(G^\epsilon(\theta))] \Rightarrow \quad \text{(good flow)}$$

$$\left. \frac{\partial}{\partial \epsilon} \right|_{\epsilon=0} \mathbb{E}_{\mathcal{P}^\epsilon(\theta)} [\phi(\theta)] = \left. \frac{\partial}{\partial \epsilon} \right|_{\epsilon=0} \mathbb{E}_{\mathcal{P}(\theta)} [\phi(G^\epsilon(\theta))] \Rightarrow \quad \text{(smoothness)}$$

Consider $\mathcal{P}^\epsilon(\theta) \propto \mathcal{P}(\theta) \exp(\epsilon h(\theta))$.

We want a flow $\theta \mapsto G^\epsilon(\theta)$ such that, for small ϵ , $G^\epsilon(\theta) \sim \mathcal{P}^\epsilon(\theta)$ and $G^0(\theta) = \theta \sim \mathcal{P}(\theta)$.

Suppose there is a scalar-valued potential function $u(\theta)$ such that

$$G^\epsilon(\theta) = \theta + \epsilon \nabla u(\theta) + O(\epsilon^2) \quad \text{as } \epsilon \rightarrow 0. \text{ Then:}$$

$$\mathbb{E}_{\mathcal{P}^\epsilon(\theta)} [\phi(\theta)] = \mathbb{E}_{\mathcal{P}(\theta)} [\phi(G^\epsilon(\theta))] \Rightarrow \quad (\text{good flow})$$

$$\left. \frac{\partial}{\partial \epsilon} \right|_{\epsilon=0} \mathbb{E}_{\mathcal{P}^\epsilon(\theta)} [\phi(\theta)] = \left. \frac{\partial}{\partial \epsilon} \right|_{\epsilon=0} \mathbb{E}_{\mathcal{P}(\theta)} [\phi(G^\epsilon(\theta))] \Rightarrow \quad (\text{smoothness})$$

$$(\star) \quad \text{Cov}_{\mathcal{P}(\theta)} (\phi(\theta), h(\theta)) = \mathbb{E}_{\mathcal{P}(\theta)} [\nabla \phi(\theta)^\top \nabla u(\theta)] \quad (\text{Connecting } h \text{ and } u)$$

Consider $\mathcal{P}^\epsilon(\theta) \propto \mathcal{P}(\theta) \exp(\epsilon h(\theta))$.

We want a flow $\theta \mapsto G^\epsilon(\theta)$ such that, for small ϵ , $G^\epsilon(\theta) \sim \mathcal{P}^\epsilon(\theta)$ and $G^0(\theta) = \theta \sim \mathcal{P}(\theta)$.

Suppose there is a scalar-valued potential function $u(\theta)$ such that

$$G^\epsilon(\theta) = \theta + \epsilon \nabla u(\theta) + O(\epsilon^2) \quad \text{as } \epsilon \rightarrow 0. \text{ Then:}$$

$$\mathbb{E}_{\mathcal{P}^\epsilon(\theta)} [\phi(\theta)] = \mathbb{E}_{\mathcal{P}(\theta)} [\phi(G^\epsilon(\theta))] \Rightarrow \quad (\text{good flow})$$

$$\left. \frac{\partial}{\partial \epsilon} \right|_{\epsilon=0} \mathbb{E}_{\mathcal{P}^\epsilon(\theta)} [\phi(\theta)] = \left. \frac{\partial}{\partial \epsilon} \right|_{\epsilon=0} \mathbb{E}_{\mathcal{P}(\theta)} [\phi(G^\epsilon(\theta))] \Rightarrow \quad (\text{smoothness})$$

$$(\star) \quad \text{Cov}_{\mathcal{P}(\theta)} (\phi(\theta), h(\theta)) = \mathbb{E}_{\mathcal{P}(\theta)} [\nabla \phi(\theta)^\top \nabla u(\theta)] \quad (\text{Connecting } h \text{ and } u)$$

Theorem: [KV86; BGL14]: Assume that $\mathbb{E}_{\mathcal{P}(\theta)} [h(\theta)^2] < \infty$ and $\mathcal{P}(\theta)$ satisfies a Poincare inequality. Then there exists a solution u to the partial differential equation

$$\mathcal{L}_{\mathcal{P}}(u) := -\nabla u^\top \nabla \log \mathcal{P} - \Delta u = h$$

such that $\|\nabla u\|_{L^2(\mathcal{P})} < \infty$ and the flow $\theta_\epsilon = \theta + \epsilon \nabla u(\theta)$ satisfies Equation $(\star)$.

Note that $\mathcal{L}_{\mathcal{P}}$ is well-known: it's the drift operator, generator of a Langevin diffusion, and a Stein operator [ABS+21; Son+21; Ros11].

How can we find $u(\theta)$ in practice?

Solving PDEs is hard. How can we easily find $u(\theta)$ in practice?

$$(\star) \quad \text{Cov}_{\mathcal{P}(\theta)}(\phi(\theta), h(\theta)) = \mathbb{E}_{\mathcal{P}(\theta)}[\nabla \phi(\theta)^\top \nabla u(\theta)] \quad (\text{Identity connecting } h \text{ and } u)$$

Basis expansion approach (the Petrov-Galerkin method [EG04; Nüs24; MM24]):

- Specify a family of real-valued target functions $\phi_n(\theta)$ for $n \in [N]$
- Specify a family of real-valued basis functions $b_m(\theta)$ for $m \in [M]$

Let $u(\theta; \beta) := \sum_{m=1}^M \beta_m b_m(\theta)$ and plug into Equation $(\star)$.

How can we find $u(\theta)$ in practice?

Solving PDEs is hard. How can we easily find $u(\theta)$ in practice?

$$(\star) \quad \text{Cov}_{\mathcal{P}(\theta)}(\phi(\theta), h(\theta)) = \mathbb{E}_{\mathcal{P}(\theta)}[\nabla \phi(\theta)^\top \nabla u(\theta)] \quad (\text{Identity connecting } h \text{ and } u)$$

Basis expansion approach (the Petrov-Galerkin method [EG04; Nüs24; MM24]):

- Specify a family of real-valued target functions $\phi_n(\theta)$ for $n \in [N]$
- Specify a family of real-valued basis functions $b_m(\theta)$ for $m \in [M]$

Let $u(\theta; \beta) := \sum_{m=1}^M \beta_m b_m(\theta)$ and plug into Equation $(\star)$.

Stacking in matrices:

$$\underbrace{\begin{bmatrix} \vdots \\ \text{Cov}_{\mathcal{P}(\theta)}(\phi_n(\theta), h(\theta)) \\ \vdots \end{bmatrix}}_{N \times 1 \text{ matrix}} = \underbrace{\begin{bmatrix} \vdots & & \\ \dots & \mathbb{E}_{\mathcal{P}(\theta)}[\nabla \phi_n(\theta)^\top \nabla b_m(\theta)] & \dots \\ \vdots & & \end{bmatrix}}_{N \times M \text{ matrix}} \begin{bmatrix} \vdots \\ \beta_m \\ \vdots \end{bmatrix}$$

This gives N linear equations in M unknowns, and the matrices can be estimated with MCMC.

Theorem [GK69; BGL14], Schwengber 2026: When $\mathcal{P}(\theta)$ is Gaussian, taking $N = M$ and $\phi_n(\theta) = b_n(\theta)$ to be the Hermite polynomials give the best finite-rank approximation to $\mathcal{L}_{\mathcal{P}}^{-1}$ in its operator norm.

Using the 1^{st} order Hermite polynomials $\phi_n(\theta) = b_n(\theta) = \theta_n$, the estimated flow is given by

$$\theta_{\epsilon} = \theta + \epsilon \underbrace{\text{Cov}_{\mathcal{P}(\theta|\mathcal{D})}(\theta, h(\theta))}_{\left. \frac{\partial \mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})}[\theta]}{\partial \epsilon} \right|_{\epsilon=0}}$$

This just moves all points in the direction of the mean shift!


An elegant flow

Theorem [GK69; BGL14], Schwengber 2026: When $\mathcal{P}(\theta)$ is Gaussian, taking $N = M$ and $\phi_n(\theta) = b_n(\theta)$ to be the Hermite polynomials give the best finite-rank approximation to $\mathcal{L}_{\mathcal{P}}^{-1}$ in its operator norm.

Using the 1st order Hermite polynomials $\phi_n(\theta) = b_n(\theta) = \theta_m$, the estimated flow is given by

$$\theta_{\epsilon} = \theta + \underbrace{\epsilon \operatorname{Cov}_{\mathcal{P}(\theta|\mathcal{D})}(\theta, h(\theta))}_{\left. \frac{\partial \mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})}[\theta]}{\partial \epsilon} \right|_{\epsilon=0}}$$

This just moves all points in the direction of the mean shift!




An elegant flow

Theorem [GK69; BGL14], Schwengber 2026: When $\mathcal{P}(\theta)$ is Gaussian, taking $N = M$ and $\phi_n(\theta) = b_n(\theta)$ to be the Hermite polynomials give the best finite-rank approximation to $\mathcal{L}_{\mathcal{P}}^{-1}$ in its operator norm.

Using the 1st order Hermite polynomials $\phi_n(\theta) = b_n(\theta) = \theta_m$, the estimated flow is given by

$$\theta_{\epsilon} = \theta + \underbrace{\epsilon \operatorname{Cov}_{\mathcal{P}(\theta|\mathcal{D})}(\theta, h(\theta))}_{\left. \frac{\partial \mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})}[\theta]}{\partial \epsilon} \right|_{\epsilon=0}}$$

This just moves all points in the direction of the mean shift!




An elegant flow

Theorem [GK69; BGL14], Schwengber 2026: When $\mathcal{P}(\theta)$ is Gaussian, taking $N = M$ and $\phi_n(\theta) = b_n(\theta)$ to be the Hermite polynomials give the best finite-rank approximation to $\mathcal{L}_{\mathcal{P}}^{-1}$ in its operator norm.

Using the 1st order Hermite polynomials $\phi_n(\theta) = b_n(\theta) = \theta_m$, the estimated flow is given by

$$\theta_\epsilon = \theta + \epsilon \underbrace{\text{Cov}_{\mathcal{P}(\theta|\mathcal{D})}(\theta, h(\theta))}_{\left. \frac{\partial \mathbb{E}_{\mathcal{P}(\theta|\mathcal{D})}[\theta]}{\partial \epsilon} \right|_{\epsilon=0}}$$

This just moves all points in the direction of the mean shift!



Suppose we have $\theta_m \sim \mathcal{P}(\theta|\mathcal{D})$, and bootstrap weights w_i .

Set $\theta_m^* = \theta_m + \frac{1}{N_S} \sum_{i=1}^{N_S} w_i \psi_i$ and compute any posterior quantities from these samples.




Figure 6: MrP Posterior Mean




Figure 6: MrP Posterior Variance




Figure 6: Covariate Shift Diagnostic

Carefully understanding the gap between IJ and the flow is ongoing work!

- Local checks of specification of Bayesian regression models
 - Use invariance to covariate shift as a diagnostic for Bayesian hierarchical models
- Frequentist variability of Bayesian posteriors with the infinitesimal jackknife
 - Use sensitivity of mean parameters to speed up bootstrapping MCMC
- Bayesian robustness with local flows
 - Use easy-to-compute approximate flows to speed up bootstrapping MCMC

Lots of exciting work to to do!

- [ABS+21] Luigi Ambrosio, Elia Brué, Daniele Semola, et al. *Lectures on optimal transport*. Vol. 130. Springer Cham, 2021.
- [BGL14] Dominique Bakry, Ivan Gentil, and Michel Ledoux. *Analysis and geometry of Markov diffusion operators*. Vol. 103. Springer, 2014.
- [Büh20] Peter Bühlmann. “Invariance, Causality and Robustness: 2018 Neyman Lecture”. In: *Statistical Science* 35.3 (2020), pp. 404–426. ISSN: 0883-4237. JSTOR: 26997912. URL: <https://www.jstor.org/stable/26997912> (visited on 08/18/2026).
- [CBR25] Rafael Cabral, David Bolin, and Håvard Rue. “Robustness, Model Checking, and Hierarchical Models”. In: *Journal of the Royal Statistical Society Series B: Statistical Methodology* 87.3 (July 14, 2025), pp. 632–652. ISSN: 1369-7412, 1467-9868. DOI: 10.1093/jrsssb/qkae107. URL: <https://academic.oup.com/jrsssb/article/87/3/632/7914952> (visited on 08/31/2026).
- [DF86] P. Diaconis and D. Freedman. “On the Consistency of Bayes Estimates”. In: *The Annals of Statistics* (1986), pp. 1–26.

- [DRM25] Antonio Di Noia, Fabrizio Ruggeri, and Antonietta Mira. “Likelihood Distortion and Bayesian Local Robustness”. In: *Bayesian Analysis* 20.4 (Dec. 1, 2025). ISSN: 1936-0975. DOI: 10.1214/25-BA1526. URL: <https://projecteuclid.org/journals/bayesian-analysis/volume-20/issue-4/Likelihood-Distortion-and-Bayesian-Local-Robustness/10.1214/25-BA1526.full> (visited on 09/01/2026).
- [EG04] Alexandre Ern and Jean-Luc Guermond. *Theory and Practice of Finite Elements*. en. Ed. by S. S. Antman, J. E. Marsden, and L. Sirovich. Vol. 159. Applied Mathematical Sciences. New York, NY: Springer, 2004. ISBN: 978-1-4419-1918-2 978-1-4757-4355-5. DOI: 10.1007/978-1-4757-4355-5. URL: <http://link.springer.com/10.1007/978-1-4757-4355-5> (visited on 01/06/2026).
- [GB26] R. Giordano and T. Broderick. “The Bayesian Infinitesimal Jackknife for Variance”. In: *arXiv preprint arXiv:2305.06466* (2026).
- [GBJ18] R. Giordano, T. Broderick, and M. Jordan. “Covariances, robustness and variational Bayes”. In: *The Journal of Machine Learning Research* 19.1 (2018), pp. 1981–2029.

- [Gel+20] A. Gelman et al. “Bayesian workflow”. In: *arXiv preprint arXiv:2011.01808* (2020).
- [Gel97] A. Gelman. “Poststratification into Many Categories Using Hierarchical Logistic Regression”. In: *Survey Methodology* 23 (1997), p. 127.
- [Gio+26] R. Giordano et al. *Locally Equivalent Weights for Multilevel Regression and Poststratification*. 2026. arXiv: 2606.04250 [stat.ME]. URL: <https://arxiv.org/abs/2606.04250>.
- [GK69] Israel Gohberg and M Kreĭn. “Introduction to the theory of linear nonselfadjoint operators in Hilbert space”. In: *Translations of mathematical monographs* (1969).
- [Gus96] P. Gustafson. “Local sensitivity of posterior expectations”. In: *The Annals of Statistics* 24.1 (1996), pp. 174–195.
- [HGV25] Geonhee Han, Andrew Gelman, and Aki Vehtari. *Efficient Scenario Analysis in Real-Time Bayesian Election Forecasting via Sequential Meta-Posterior Sampling*. arXiv.org. Oct. 21, 2025. URL: <https://arxiv.org/abs/2510.19133v1> (visited on 05/27/2026).
- [HM23] J. Huggins and J. Miller. “Reproducible model selection using bagged posteriors”. In: *Bayesian Analysis* 18.1 (2023), pp. 79–104.

- [JLR25] Feng Ji, JoonHo Lee, and Sophia Rabe-Hesketh. “Valid Standard Errors for Bayesian Quantile Regression With Clustered and Independent Data”. In: *Journal of Educational and Behavioral Statistics* (2025), p. 10769986251379738. ISSN: 1076-9986. DOI: 10.3102/10769986251379738. URL: <https://doi.org/10.3102/10769986251379738> (visited on 08/31/2026).
- [Kuh+24] Swen Kuh et al. “Using Leave-One-out Cross Validation (LOO) in a Multilevel Regression and Poststratification (MRP) Workflow: A Cautionary Tale”. In: *Statistics in Medicine* 43.5 (2024), pp. 953–982. ISSN: 1097-0258. DOI: 10.1002/sim.9964. URL: <https://onlinelibrary.wiley.com/doi/abs/10.1002/sim.9964> (visited on 05/18/2026).
- [KV86] C. Kipnis and S. R. S. Varadhan. “Central Limit Theorem for Additive Functionals of Reversible Markov Processes and Applications to Simple Exclusions”. In: *Communications in Mathematical Physics* 104.1 (Mar. 1986), pp. 1–19. ISSN: 0010-3616, 1432-0916. DOI: 10.1007/BF01210789. URL: <http://link.springer.com/10.1007/BF01210789> (visited on 09/03/2026).

- [LJ26] Nanyu Luo and Feng Ji. *Robust Standard Errors for Bayesian Posterior Functionals via the Infinitesimal Jackknife*. 2026. DOI: 10.48550/arXiv.2604.03398. arXiv: 2604.03398 [stat.ME]. URL: <http://arxiv.org/abs/2604.03398> (visited on 07/24/2026). Pre-published.
- [MJK20] Krikamol Muandet, Wittawat Jitkrittum, and Jonas Kübler. “Kernel Conditional Moment Test via Maximum Moment Restriction”. In: *Proceedings of the 36th Conference on Uncertainty in Artificial Intelligence (UAI)*. Conference on Uncertainty in Artificial Intelligence. PMLR, Aug. 27, 2020, pp. 41–50. URL: <https://proceedings.mlr.press/v124/muandet20a.html> (visited on 08/31/2026).
- [MM24] Aimee Maurais and Youssef Marzouk. *Sampling in Unit Time with Kernel Fisher-Rao Flow*. June 5, 2024. DOI: 10.48550/arXiv.2401.03892. arXiv: 2401.03892 [stat.CO]. URL: <http://arxiv.org/abs/2401.03892> (visited on 09/03/2026). Pre-published.
- [Nüs24] Nikolas Nüsken. *Stein Transport for Bayesian Inference*. Nov. 28, 2024. DOI: 10.48550/arXiv.2409.01464. arXiv: 2409.01464 [stat.ML]. URL: <http://arxiv.org/abs/2409.01464> (visited on 09/03/2026). Pre-published.

- [Pap+21] G. Papamakarios et al. “Normalizing flows for probabilistic modeling and inference”. In: *Journal of Machine Learning Research* 22.57 (2021), pp. 1–64.
- [Roj+] Mateo Rojas-Carulla et al. “Invariant Models for Causal Transfer Learning”. In: ().
- [Ros11] Nathan Ross. “Fundamentals of Stein’s method”. In: (2011).
- [Shi00] Hidetoshi Shimodaira. “Improving Predictive Inference under Covariate Shift by Weighting the Log-Likelihood Function”. In: *Journal of Statistical Planning and Inference* 90.2 (Oct. 2000), pp. 227–244. ISSN: 03783758. DOI: 10.1016/S0378-3758(00)00115-4. URL: <https://linkinghub.elsevier.com/retrieve/pii/S0378375800001154> (visited on 09/03/2026).
- [Son+21] Yang Song et al. *Score-Based Generative Modeling through Stochastic Differential Equations*. 2021. arXiv: 2011.13456 [cs.LG]. URL: <https://arxiv.org/abs/2011.13456>.

- [VGG17] Aki Vehtari, Andrew Gelman, and Jonah Gabry. “Practical Bayesian Model Evaluation Using Leave-One-out Cross-Validation and WAIC”. In: *Statistics and Computing* 27.5 (Sept. 1, 2017), pp. 1413–1432. ISSN: 1573-1375. DOI: 10.1007/s11222-016-9696-4. URL: <https://doi.org/10.1007/s11222-016-9696-4> (visited on 05/27/2026).
- [von47] R. von Mises. “On the asymptotic distribution of differentiable statistical functions”. In: *The Annals of Mathematical Statistics* 18.3 (1947), pp. 309–348.
- [WDH26] Yu Wang, Jie Ding, and Jonathan H. Huggins. *Accurate Large-sample Uncertainty Quantification Using Stochastic Gradient Markov Chain Monte Carlo*. 2026. DOI: 10.48550/arXiv.2606.00293. arXiv: 2606.00293 [cs.LG]. URL: <http://arxiv.org/abs/2606.00293> (visited on 08/31/2026). Pre-published.
- [YB24] B. Yu and R. Barter. *Veridical Data Science: The Practice of Responsible Data Analysis and Decision Making*. MIT Press, 2024.