

Variational Methods for Latent Variable Problems

Ryan Giordano (for Johns Hopkins Biostats BLAST working group)

Oct, 2021

Massachusetts Institute of Technology

Outline for today:

- Some examples of latent variable models
- A template: The Neyman-Scott “paradox” and marginalization
- Bayesian versus frequentist approaches to marginalization
- The classical EM algorithm (in brief)


Next week, we will build on these ideas to present more general variational inference.

KL divergence exercises

$p(\theta)$ = Correlated bivariate normal

$\mathcal{Q} = \{\text{All bivariate normals}\}$

What is $q^*(\theta) = \operatorname{argmin}_{q \in \mathcal{Q}} \text{KL}(q(\theta) || p(\theta))$?




KL divergence exercises

$p(\theta)$ = Correlated bivariate normal

$\mathcal{Q} = \{\text{All bivariate normals}\}$

What is $q^*(\theta) = \underset{q \in \mathcal{Q}}{\operatorname{argmin}} \operatorname{KL}(q(\theta) || p(\theta))$?




KL divergence exercises

$p(\theta)$ = Correlated bivariate normal

$\mathcal{Q} = \{\text{Independent bivariate normals}\}$

What is $q^*(\theta) = \operatorname{argmin}_{q \in \mathcal{Q}} \text{KL}(q(\theta) || p(\theta))$?




KL divergence exercises

$p(\theta)$ = Correlated bivariate normal

$\mathcal{Q} = \{\text{Independent bivariate normals}\}$

What is $q^*(\theta) = \operatorname{argmin}_{q \in \mathcal{Q}} \text{KL}(q(\theta) || p(\theta))$?




KL divergence exercises

$p(\theta)$ = Correlated bivariate normal

$\mathcal{Q} = \{\text{Independent bivariate normals}\}$

What is $q^*(\theta) = \operatorname{argmin}_{q \in \mathcal{Q}} \text{KL}(p(\theta) || q(\theta))$?




KL divergence exercises

$p(\theta)$ = Correlated bivariate normal

$\mathcal{Q} = \{\text{Independent bivariate normals}\}$

What is $q^*(\theta) = \operatorname{argmin}_{q \in \mathcal{Q}} \text{KL}(p(\theta) || q(\theta))$?



KL divergence exercises

Recall that


$$\text{KL}(q(\theta) || p(\theta)) = - \mathbb{E}_{q(\theta)} [\log p(\theta)] + \mathbb{E}_{q(\theta)} [\log q(\theta)]$$

$p(\theta)$ = Correlated bivariate normal

$$\mathcal{Q} = \{\text{Bivariate normals}\}$$

What is $q^*(\theta) =$

$$\underset{q \in \mathcal{Q}}{\text{argmin}} \left(- \mathbb{E}_{q(\theta)} [\log p(\theta)] + \mathbb{E}_{q(\theta)} [\log q(\theta)] \right)?$$



KL divergence exercises

Recall that


$$\text{KL}(q(\theta) || p(\theta)) = - \mathbb{E}_{q(\theta)} [\log p(\theta)] + \mathbb{E}_{q(\theta)} [\log q(\theta)]$$

$p(\theta)$ = Correlated bivariate normal

$$\mathcal{Q} = \{\text{Bivariate normals}\}$$

What is $q^*(\theta)$ =

$$\underset{q \in \mathcal{Q}}{\text{argmin}} \left(- \mathbb{E}_{q(\theta)} [\log p(\theta)] + \mathbb{E}_{q(\theta)} [\log q(\theta)] \right)?$$



KL divergence exercises

Recall that


$$\text{KL}(q(\theta) || p(\theta)) = - \mathbb{E}_{q(\theta)} [\log p(\theta)] + \mathbb{E}_{q(\theta)} [\log q(\theta)]$$

$p(\theta)$ = Correlated bivariate normal

$$\mathcal{Q} = \{\text{Bivariate normals}\}$$

What is $q^*(\theta) =$

$$\underset{q \in \mathcal{Q}}{\text{argmin}} \left(- \mathbb{E}_{q(\theta)} [\log p(\theta)] + \mathbb{E}_{q(\theta)} [\log q(\theta)] \right)?$$



KL divergence exercises

Recall that


$$\text{KL}(q(\theta) || p(\theta)) = -\mathbb{E}_{q(\theta)} [\log p(\theta)] + \mathbb{E}_{q(\theta)} [\log q(\theta)]$$

$p(\theta)$ = Correlated bivariate normal

$$\mathcal{Q} = \{\text{Bivariate normals}\}$$

What is $q^*(\theta)$ =

$$\underset{q \in \mathcal{Q}}{\text{argmin}} \left(-\mathbb{E}_{q(\theta)} [\log p(\theta)] + \mathbb{E}_{q(\theta)} [\log q(\theta)] \right)?$$




KL divergence exercises

$p(\theta)$ = Correlated bivariate normal

$\mathcal{Q} = \{\text{Point masses}\}$

What is $q^*(\theta) = \operatorname{argmin}_{q \in \mathcal{Q}} \text{KL}(q(\theta) || p(\theta))$?




KL divergence exercises

$p(\theta)$ = Correlated bivariate normal

$\mathcal{Q} = \{\text{Point masses}\}$

What is $q^*(\theta) = \operatorname{argmin}_{q \in \mathcal{Q}} \text{KL}(q(\theta) || p(\theta))$?




KL divergence exercises

$p(\theta)$ = Correlated bivariate normal

$\mathcal{Q} = \{\text{BVN with small, fixed variance}\}$

What is $q^*(\theta) = \underset{q \in \mathcal{Q}}{\operatorname{argmin}} \operatorname{KL}(q(\theta) || p(\theta))$?



KL divergence exercises

$p(\theta)$ = Correlated bivariate normal

$\mathcal{Q} = \{\text{BVN with small, fixed variance}\}$

What is $q^*(\theta) = \underset{q \in \mathcal{Q}}{\operatorname{argmin}} \operatorname{KL}(q(\theta) || p(\theta))$?

