

Regents Junior Faculty Fellowship

Ryan Giordano

I propose to spend the summer of 2024 working on two collaborative research projects. The first, “neural network classifiers for Bayesian posteriors,” promises to introduce a completely new set of Bayesian inference techniques with different computational tradeoffs than existing methods. The second, “black-box computable diagnostic weights for survey sampling,” will bring a much-needed set of diagnostic tools to the vast majority of modern applied survey sampling. These two projects are different in scope — the first represents ground-breaking methodological research, and the second an application of my existing research to an urgent applied problem — but each rests on and contributes to my existing work on approximate Bayesian computation and sensitivity analysis.

Neural network classifiers for Bayesian posteriors

Bayesian statistical techniques are a conceptually powerful set of tools for representing and quantifying uncertainty, and are increasingly popular across the physical and social sciences. Often, a statistical analysis involves a single quantity of interest, such as the effect of a policy intervention [Meager, 2019], the type of an astronomical object [Regier et al., 2019], the outcome of an election [Gelman and Heidemanns, 2020], or the identity of an ancestral genetic population [Pritchard et al., 2000], and Bayesian statistics able to propagate uncertainty from a any unknown latent modeling quantities to the final estimate. But this conceptual strength is a computational weakness, since even approximately accounting for a large number of latent quantities is computationally intensive. Bayesian estimates often take hours to days to compute, and it is of considerable interest to develop computationally efficient, approximate Bayesian procedures [Blei et al., 2017, AABI, 2024].

Together with a staff scientist at LBNL, I have recently developed a new approach to Bayesian inference based on neural network classifiers (NNC). The idea is derived from a technique for point estimation in simulation-based inference (SBI) (i.e., inference in problems without a tractable likelihood function), a technique I will refer to as SBI-NNC [Cranmer et al., 2020]. Rather than learning a distribution directly (as is done in nearly all other ML-based Bayesian learning, e.g. [Papamakarios et al., 2021]), this technique uses the fact that optimal classifiers learn likelihood ratios. I have shown that a variant of the SBI-NNC trick can be applied to learn Bayesian marginals without having to learn the distribution of all the latent variables, at the cost of training a NNC on a single classification task. I will refer to this technique as Bayes-NNC.

Both classical Bayesian procedures and the existing SBI-NNC trick are difficult to validate in practice, due to the lack of a computable ground truth. Amazingly, Bayes-NNC does not suffer from this shortcoming, and its accuracy is readily testable using simulation-based calibration (SBC) [Talts et al., 2018]. Put together, Bayes-NNC and SBC offer a way to learn Bayesian posterior densities with strong, computable statistical guarantees. Interestingly, though well-known, SBC is rarely used in practice, since it is typically computationally prohibitive to compute the posterior at many different datapoints. However, Bayes-NNC learns the posterior for many datasets simultaneously, permitting efficient use of SBC in practice.

To my knowledge, there are no existing Bayesian techniques that offer the advantages of Bayes-NNC and SBC. Bayesian approaches to simulation-based inference are not new, but existing techniques are built on high-dimensional density approximation, such as normalizing flows [Cranmer et al., 2020]. As with other approximate inference techniques, this set of tools

approximates the entire posterior, even when only a low-dimensional marginal is of interest. To the best of my (and my LBNL collaborator’s) knowledge, the technique above is new, and offers a distinct set of computational tradeoffs relative to existing methods.

I would stress that Bayes–NNC is not at all limited to SBI problems. There are many problems with tractable likelihoods which are easy to simulate from and extremely computationally intensive to sample from. There is every reason to believe that training a NNC will be competitive with existing computational techniques, even neglecting the advantage of accompanying statistical guarantees. A large number of fundamental statistical applications, superficially distinct from SBI problems, are immediate candidates for Bayes–NNC, such as sparse model selection (e.g. [Ročková and George, 2018]), Gaussian process hyperparameter calibration (e.g. [Hensman et al., 2015]), and unsupervised clustering (e.g. McAuliffe et al. [2006]).

Black–box computable diagnostic weights for survey sampling

Most modern surveys — such as polling about the upcoming presidential election — must overcome the fact that their sampled population is different from the target population. For example, the set of people responding to an internet survey about political preferences is likely to differ systematically from the full population of voters, and it is extremely useful to be able to check that the re-weighting is accurate, for example by checking that key demographic variables are balanced by the re-weighting. Unfortunately, the most accurate and most commonly used statistical procedures for inferring the polling responses of rare demographic groups are nonlinear, and so do not readily admit diagnostic weights.

In collaboration with a UC Berkeley professor of public policy, I have shown that one can compute “local diagnostic weights” for non-linear statistical procedure, provided a much-needed diagnostic that is currently unavailable. The local weights I derive are closely related to the classical “influence function” of robust statistics. Though the influence function is well-studied in the frequentist literature, it has been relatively neglected in the Bayesian literature (with my own recent work being an exception).

Importantly, the local weights can be automatically computed with a small library built on top of existing open-source software which is commonly used for survey analysis (CITE). I have already implemented a similar package for a different style of sensitivity analysis (CITE), and we hope to be able to release open-source software relatively quickly.

References

- AABI. 6th symposium on advances in approximate Bayesian inference, 2024. URL <https://approximateinference.org/>.
- D. Blei, A. Kucukelbir, and J. McAuliffe. Variational inference: A review for statisticians. *Journal of the American Statistical Association*, 112(518):859–877, 2017.
- K. Cranmer, J. Brehmer, and G. Louppe. The frontier of simulation-based inference. *Proceedings of the National Academy of Sciences*, 117(48):30055–30062, 2020.
- A. Gelman and M. Heidemanns. The Economist: Forecasting the US elections., 2020. URL <https://projects.economist.com/us-2020-forecast/president>. Data and model accessed Oct., 2020.
- J. Hensman, A. Matthews, and Z. Ghahramani. Scalable variational Gaussian process classification. In *Artificial Intelligence and Statistics*, pages 351–360. PMLR, 2015.
- J. McAuliffe, D. Blei, and M. Jordan. Nonparametric empirical bayes for the Dirichlet process mixture model. *Statistics and Computing*, 16:5–14, 2006.
- R. Meager. Understanding the average impact of microcredit expansions: A Bayesian hierarchical analysis of seven randomized experiments. *American Economic Journal: Applied Economics*, 11(1):57–91, 2019.
- G. Papamakarios, E. Nalisnick, D. Rezende, S. Mohamed, and B. Lakshminarayanan. Normalizing flows for probabilistic modeling and inference. *Journal of Machine Learning Research*, 22(57):1–64, 2021.
- J. Pritchard, M. Stephens, and P. Donnelly. Inference of population structure using multilocus genotype data. *Genetics*, 155(2):945–959, 2000.
- J. Regier, K. Fischer, K. Pamnany, A. Noack, J. Revels, M. Lam, S. Howard, R. Giordano, D. Schlegel, and J. McAuliffe. Cataloging the visible universe through Bayesian inference in Julia at petascale. *Journal of Parallel and Distributed Computing*, 127:89–104, 2019.
- V. Ročková and E. George. The spike-and-slab lasso. *Journal of the American Statistical Association*, 113(521):431–444, 2018.
- S. Talts, M. Betancourt, D. Simpson, A. Vehtari, and A. Gelman. Validating Bayesian inference algorithms with simulation-based calibration. *arXiv preprint arXiv:1804.06788*, 2018.